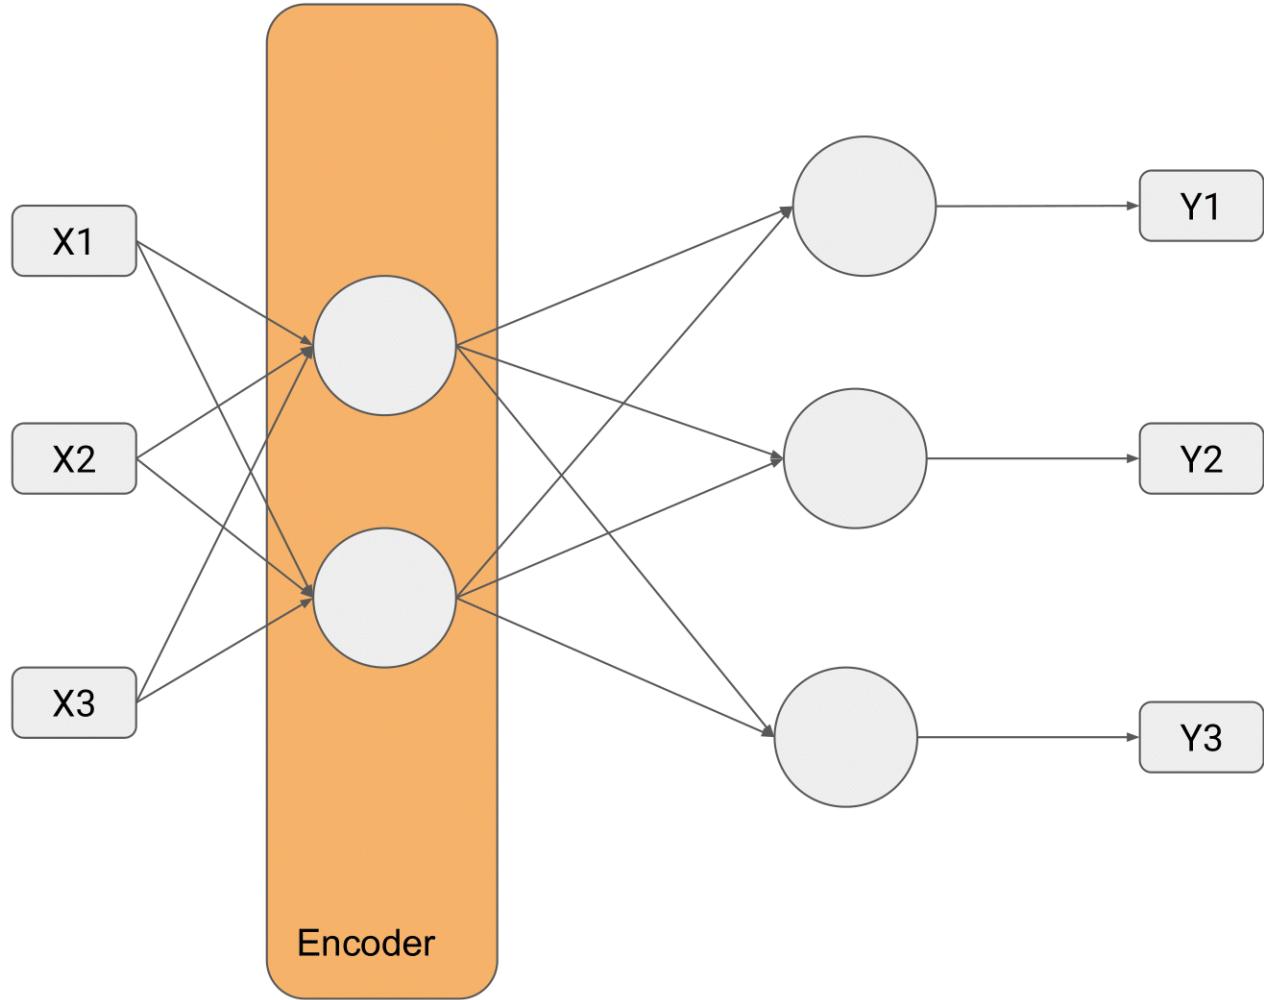
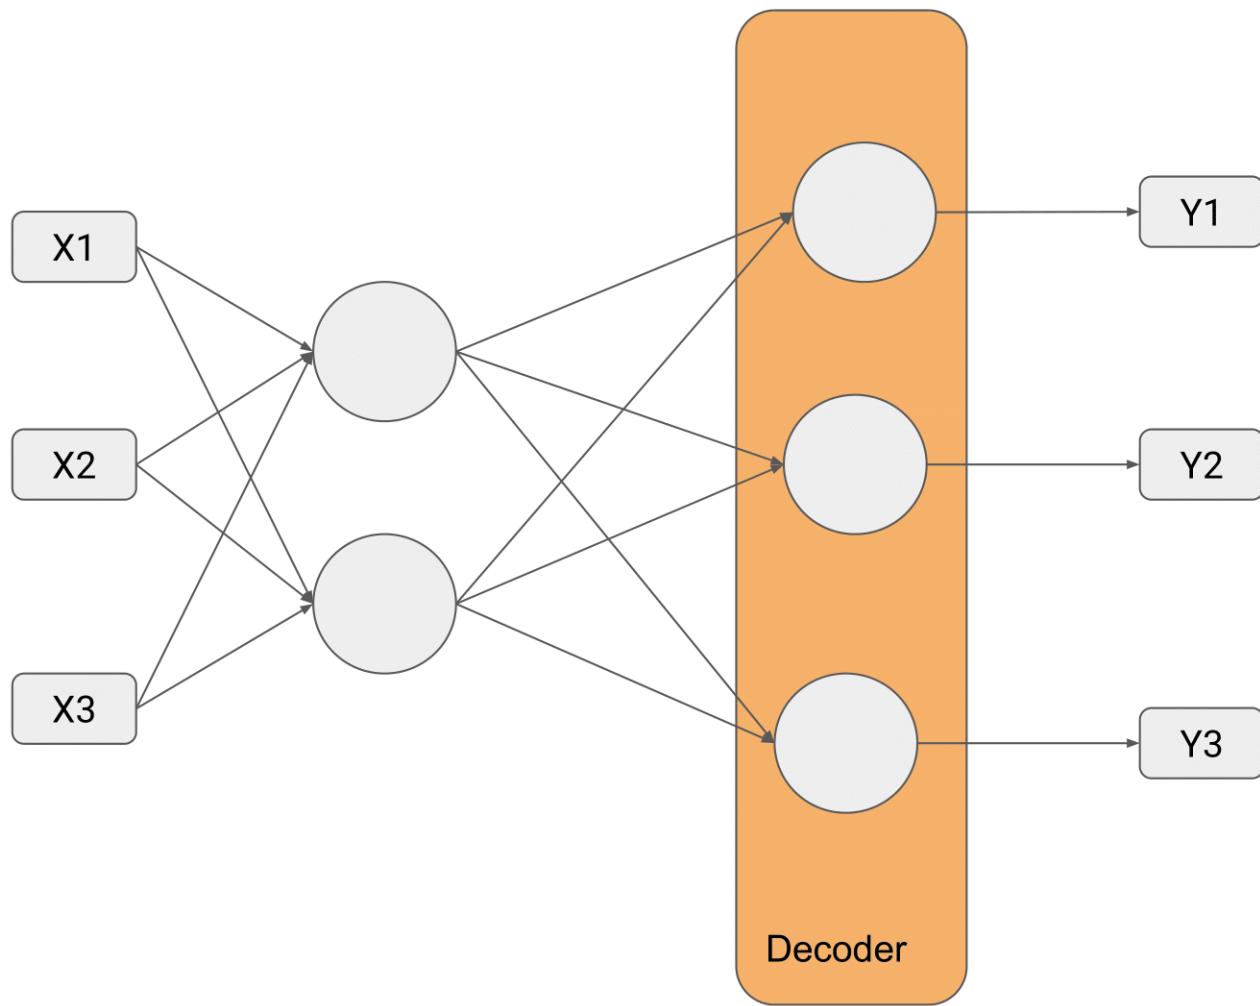
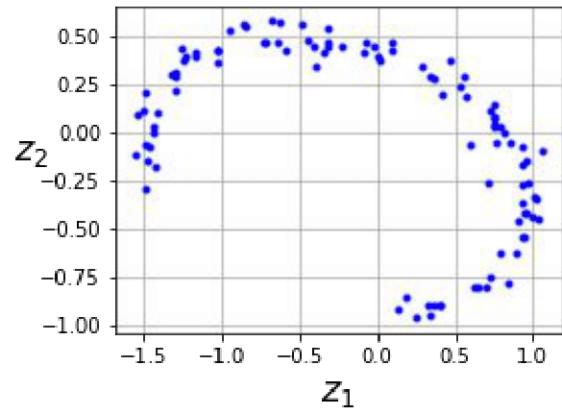
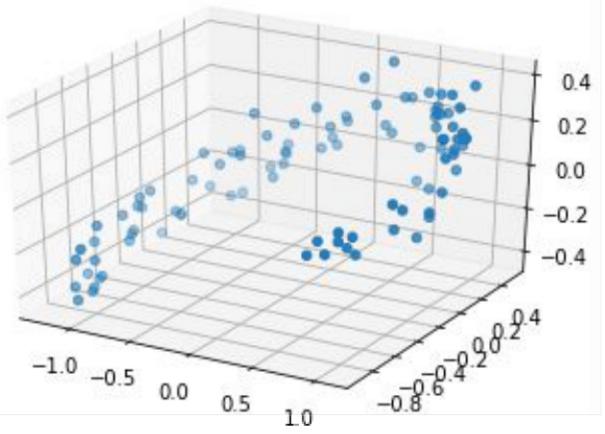


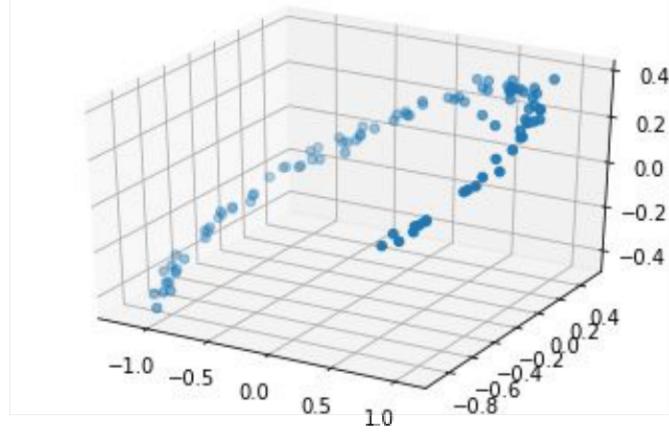
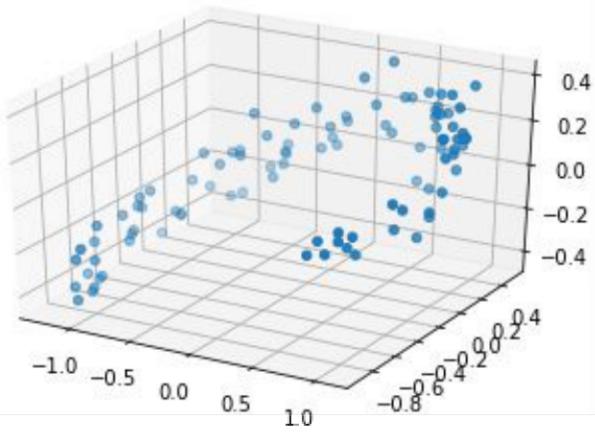
What are AutoEncoders?

- Neural networks capable of learning dense representations of input data without supervision
 - Training data is not labelled
- Useful for dimensionality reduction and for visualization
- Can be used to generate new data that resembles input data
- In practice they
 - Copy input to output
 - They learn efficient ways to represent data

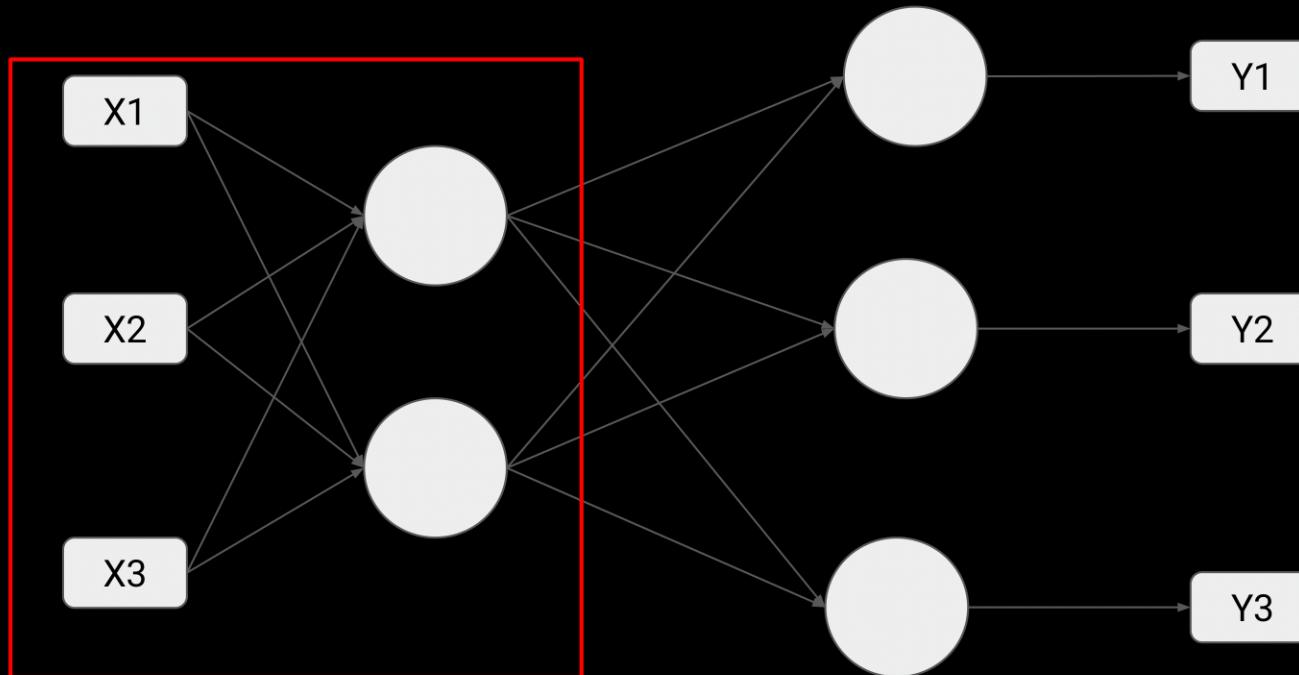








```
encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])  
  
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])  
  
autoencoder = keras.models.Sequential([encoder, decoder])  
  
autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=1.5))
```

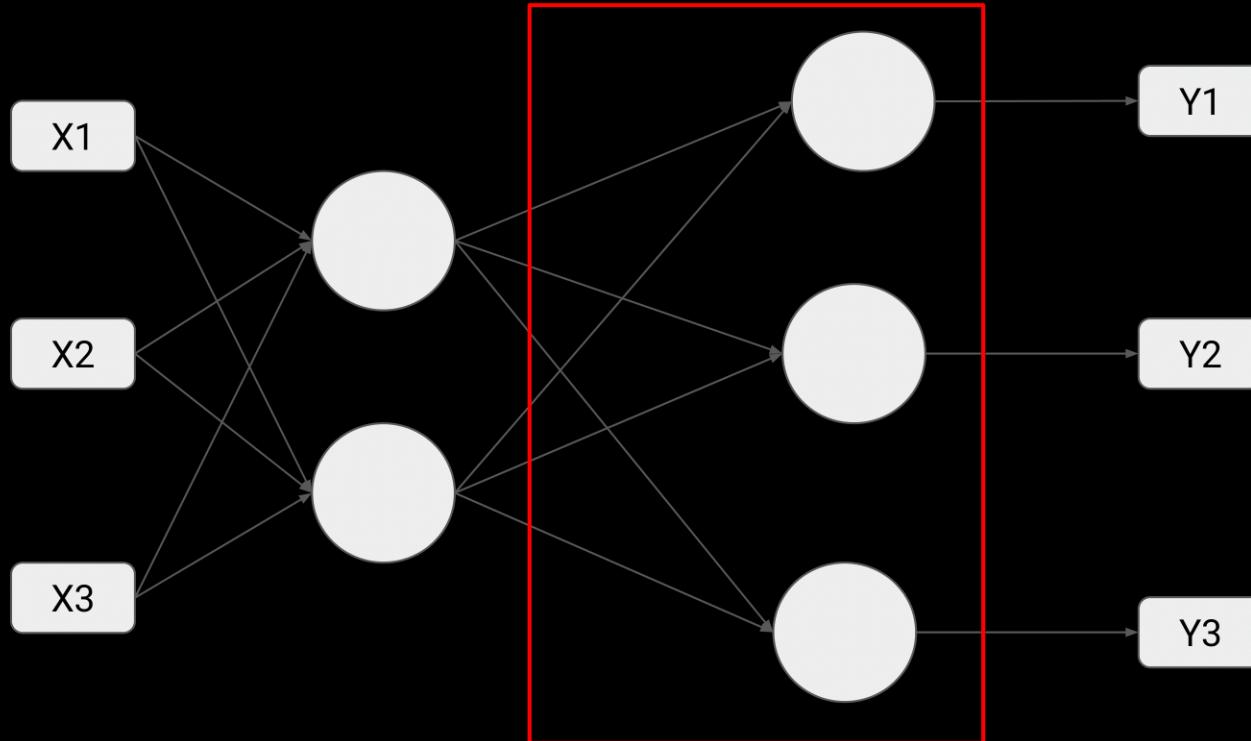


```
encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])
```

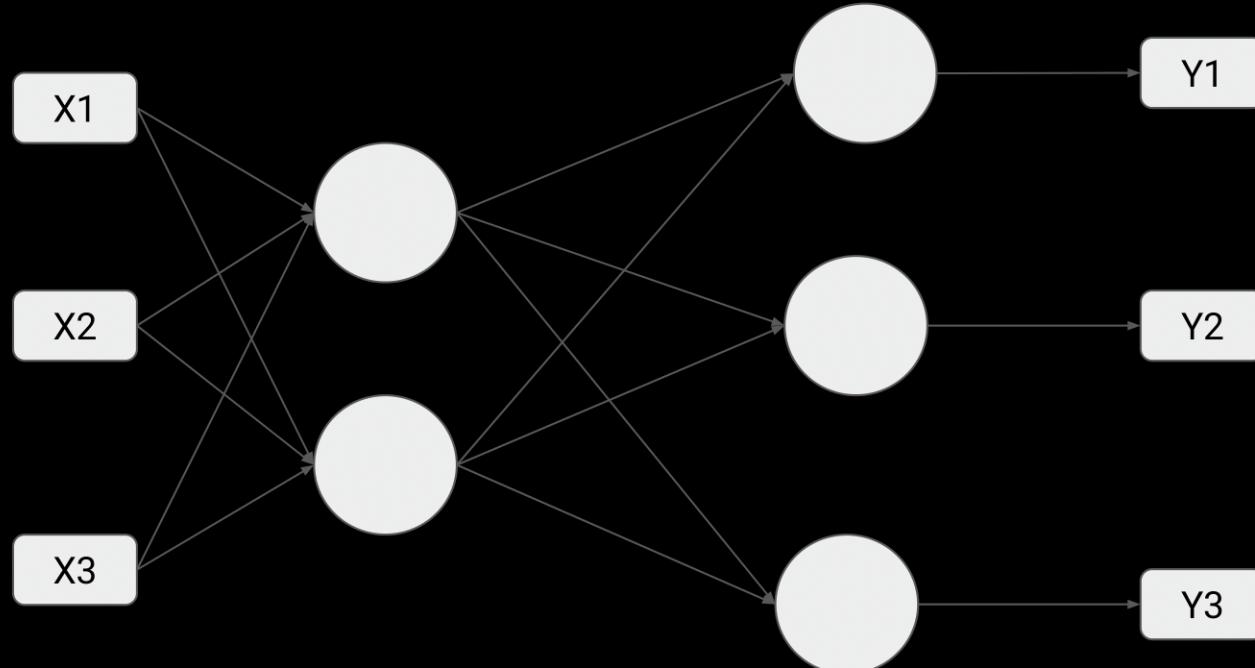
```
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])
```

```
autoencoder = keras.models.Sequential([encoder, decoder])
```

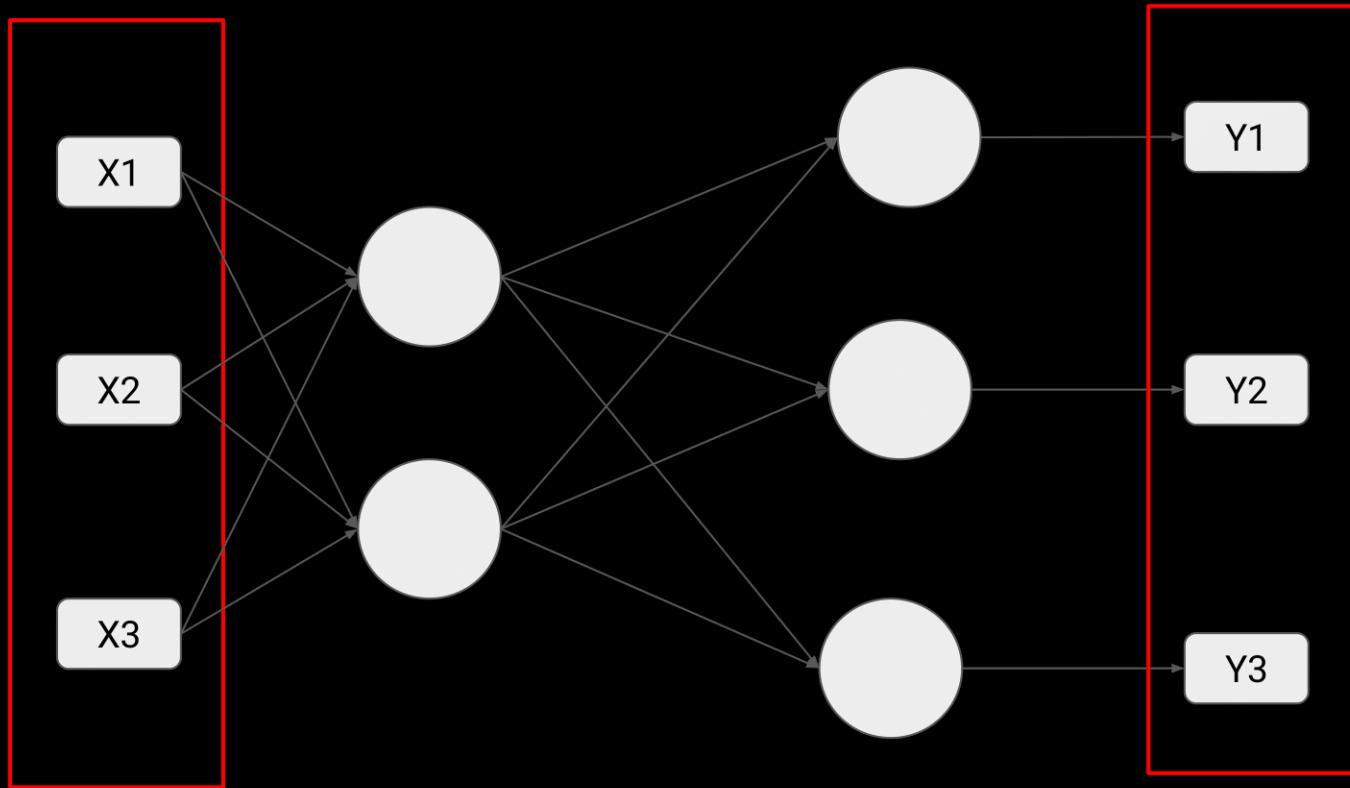
```
autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=1.5))
```



```
encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])  
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])  
  
autoencoder = keras.models.Sequential([encoder, decoder])  
  
autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=1.5))
```



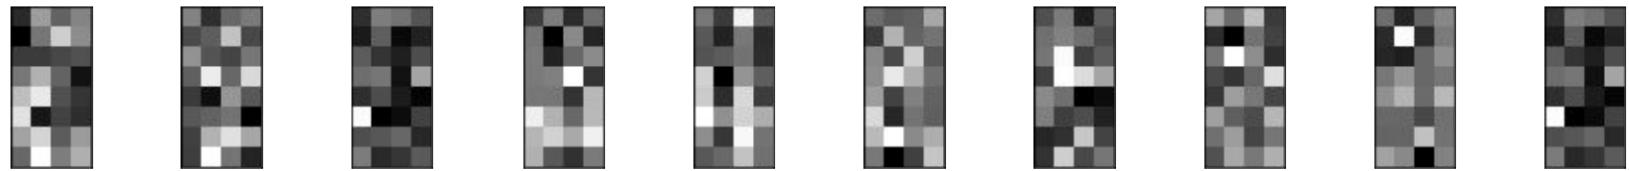
```
history = autoencoder.fit(X_train, X_train, epochs=200)
```



```
codings = encoder.predict(data)
```

```
decodings = decoder.predict(codings)
```

5 3 8 8 7 5 1 1 9 8



5 3 8 8 7 5 1 1 9 8

```
inputs = tf.keras.layers.Input(shape=(784,))

def simple_autoencoder():
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(inputs)
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(encoder)
    return encoder, decoder

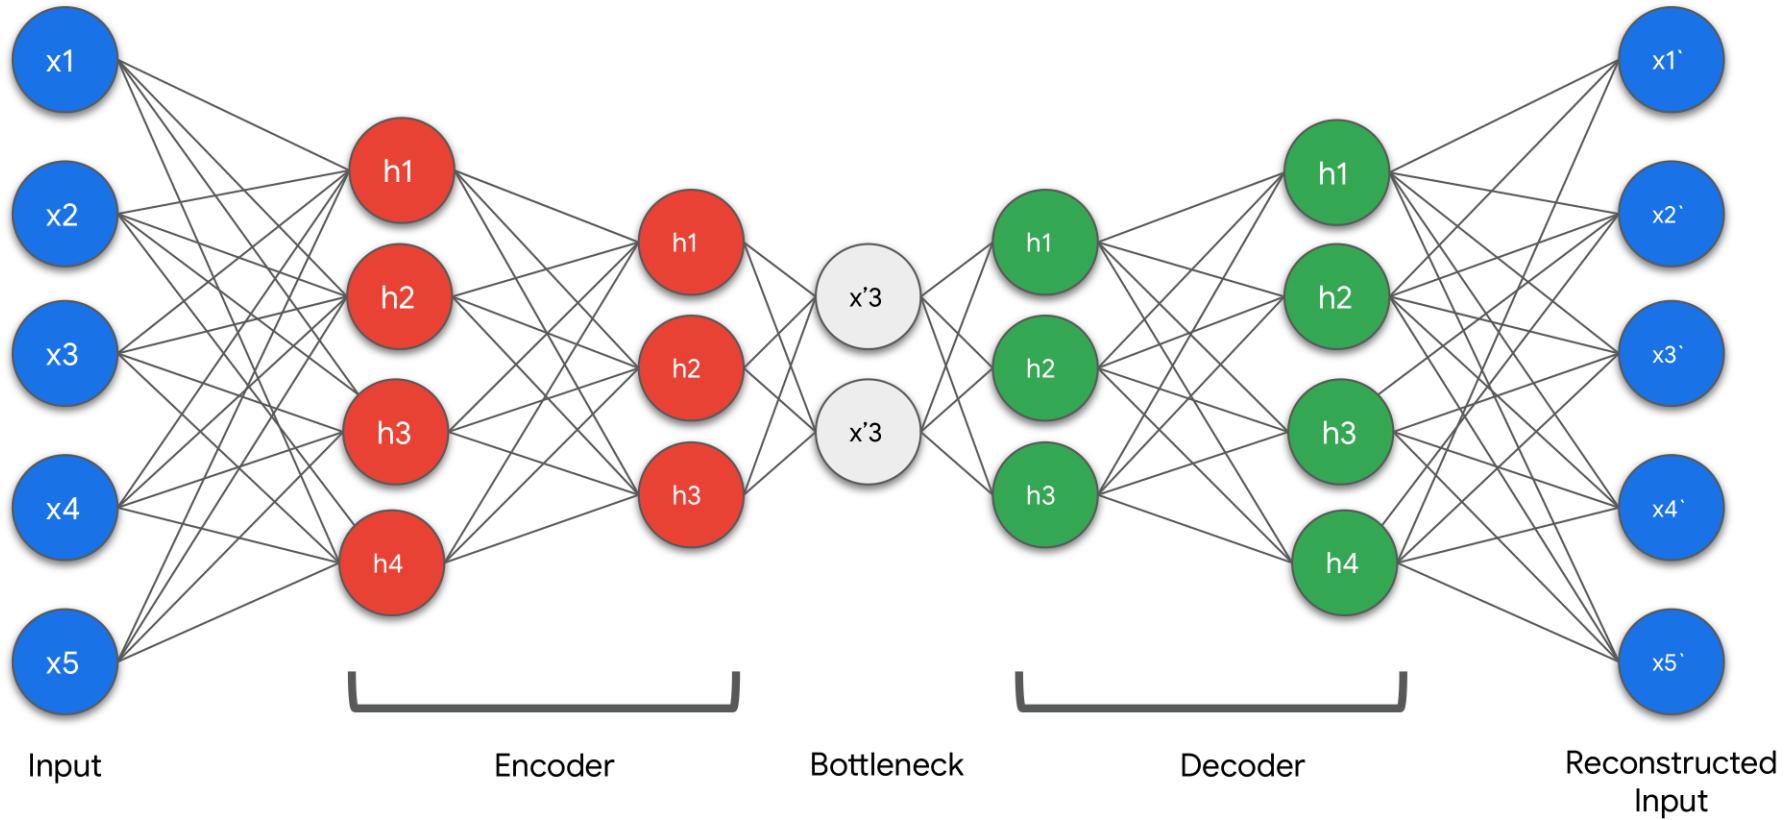
encoder_output, decoder_output = simple_autoencoder()

encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_output)

autoencoder_model = tf.keras.Model(inputs=inputs, outputs=decoder_output)
```

```
autoencoder_model.compile(  
    optimizer=tf.keras.optimizers.Adam(),  
    loss='binary_crossentropy')
```

Stacked Auto-Encoders



```
inputs = tf.keras.layers.Input(shape=(784,))

def deep_autoencoder():
    encoder = tf.keras.layers.Dense(units=128, activation='relu')(inputs)
    encoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)
    encoder = tf.keras.layers.Dense(units=32, activation='relu')(encoder)

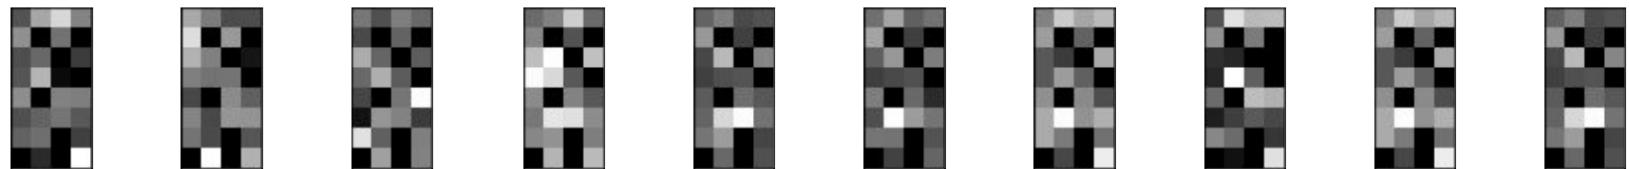
    decoder = tf.keras.layers.Dense(units=64, activation='relu')(encoder)
    decoder = tf.keras.layers.Dense(units=128, activation='relu')(decoder)
    decoder = tf.keras.layers.Dense(units=784, activation='sigmoid')(decoder)

    return encoder, decoder

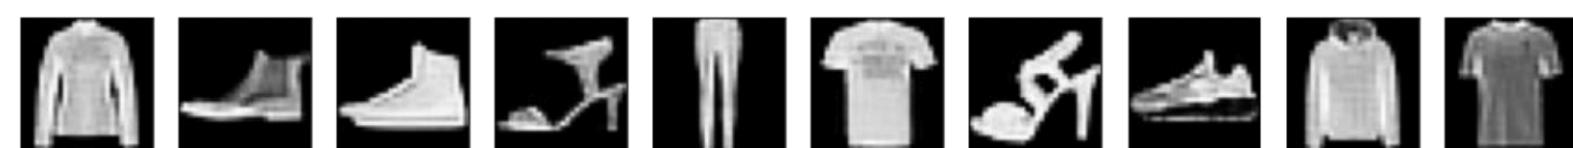
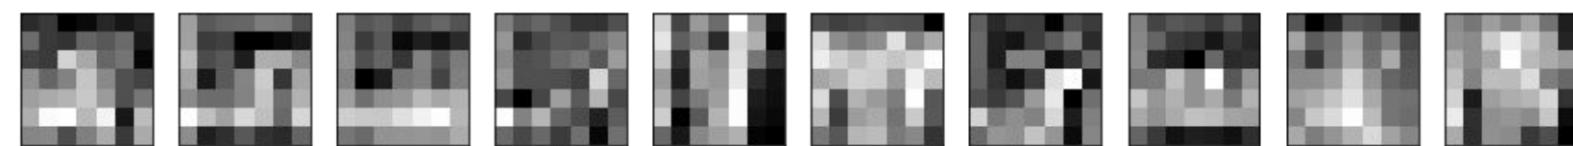
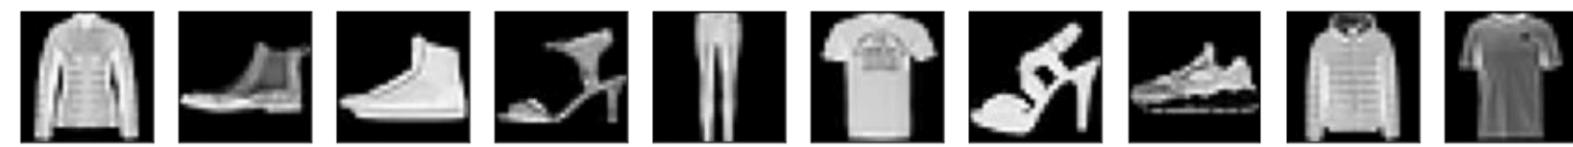
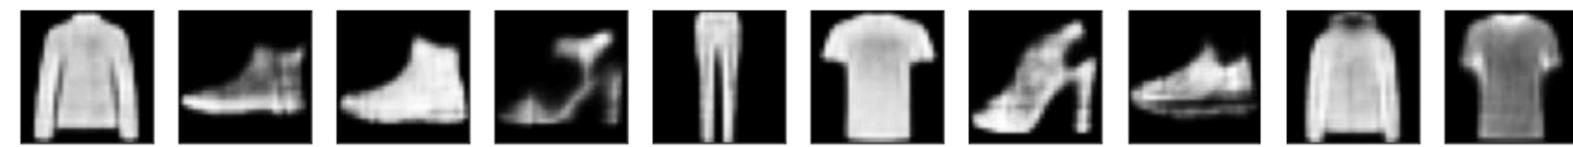
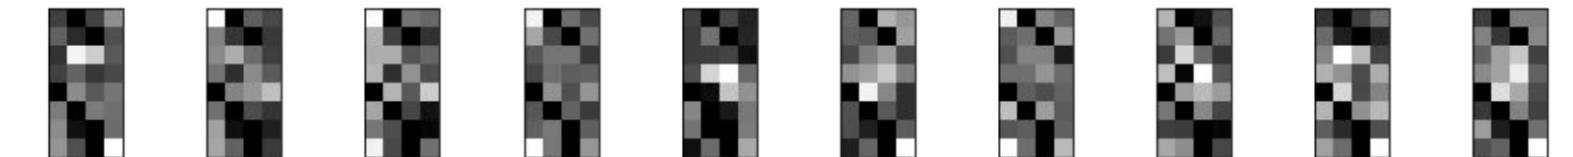
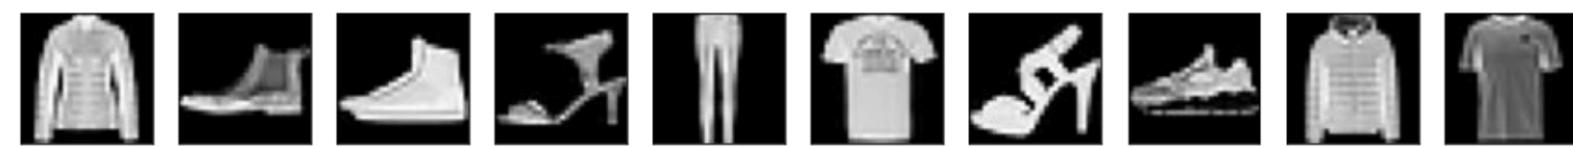
deep_encoder_output, deep_autoencoder_output = deep_autoencoder()

deep_encoder_model = tf.keras.Model(inputs=inputs, outputs=deep_encoder_output)
deep_autoencoder_model = tf.keras.Model(inputs=inputs, outputs=deep_autoencoder_output)
```

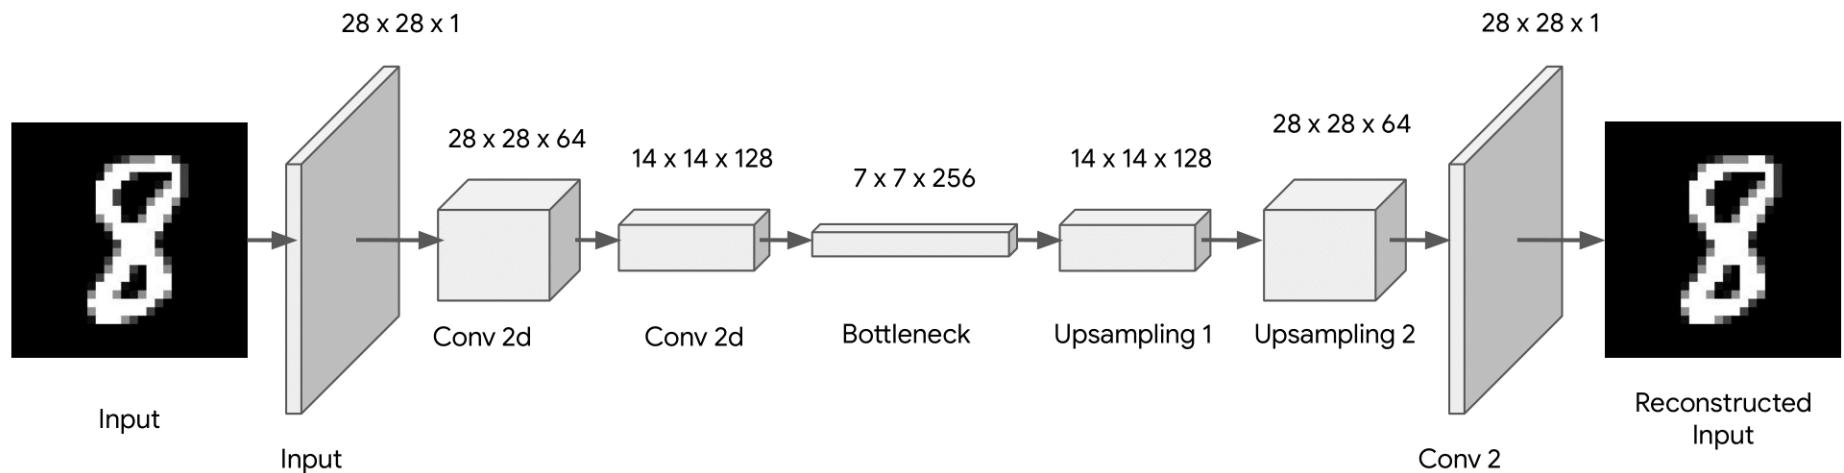
1 4 5 8 6 6 6 1 6 6



1 4 5 8 6 6 6 1 6 6



Convolutional Auto-Encoders



```
def encoder(inputs):

    conv_1 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),
                                    activation='relu', padding='same')(inputs)

    max_pool_1 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_1)

    conv_2 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),
                                    activation='relu', padding='same')(max_pool_1)

    max_pool_2 = tf.keras.layers.MaxPooling2D(pool_size=(2,2))(conv_2)

    return max_pool_2
```

```
def bottle_neck(inputs):
    bottle_neck = tf.keras.layers.Conv2D(filters=256, kernel_size=(3,3),
                                         activation='relu', padding='same')(inputs)

    encoder_visualization = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),
                                                   activation='sigmoid',
                                                   padding='same')(bottle_neck)

    return bottle_neck, encoder_visualization
```

```
def decoder(inputs):
    conv_1 = tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3),
                                  activation='relu', padding='same')(inputs)
    up_sample_1 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_1)

    conv_2 = tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3),
                                  activation='relu', padding='same')(up_sample_1)
    up_sample_2 = tf.keras.layers.UpSampling2D(size=(2,2))(conv_2)

    conv_3 = tf.keras.layers.Conv2D(filters=1, kernel_size=(3,3),
                                  activation='sigmoid',
                                  padding='same')(up_sample_2)

    return conv_3
```

```
def convolutional_auto_encoder():
    inputs = tf.keras.layers.Input(shape=(28, 28, 1,))

    encoder_output = encoder(inputs)

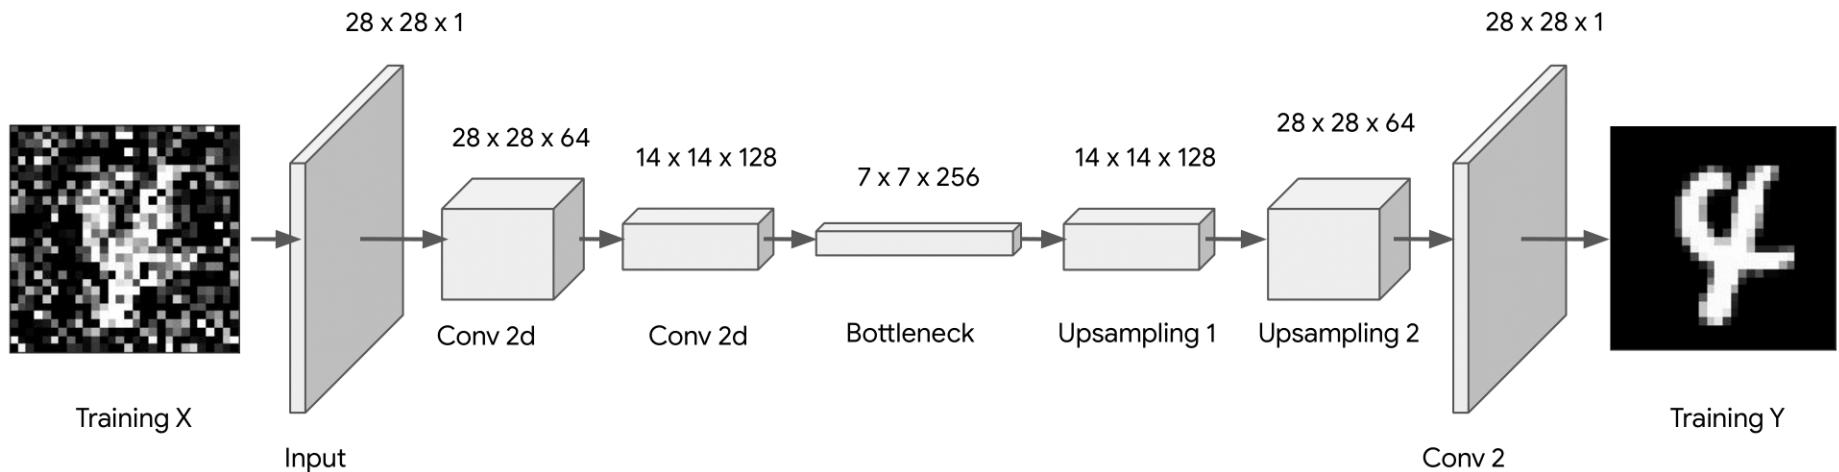
    bottleneck_output, encoder_visualization = bottle_neck(encoder_output)

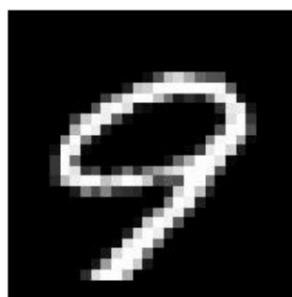
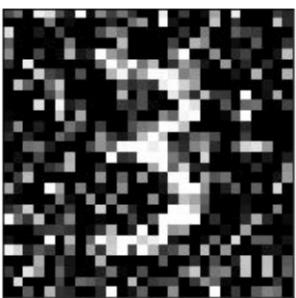
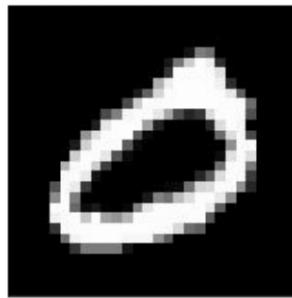
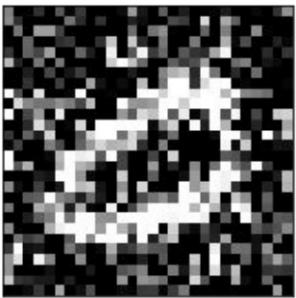
    decoder_output = decoder(bottleneck_output)

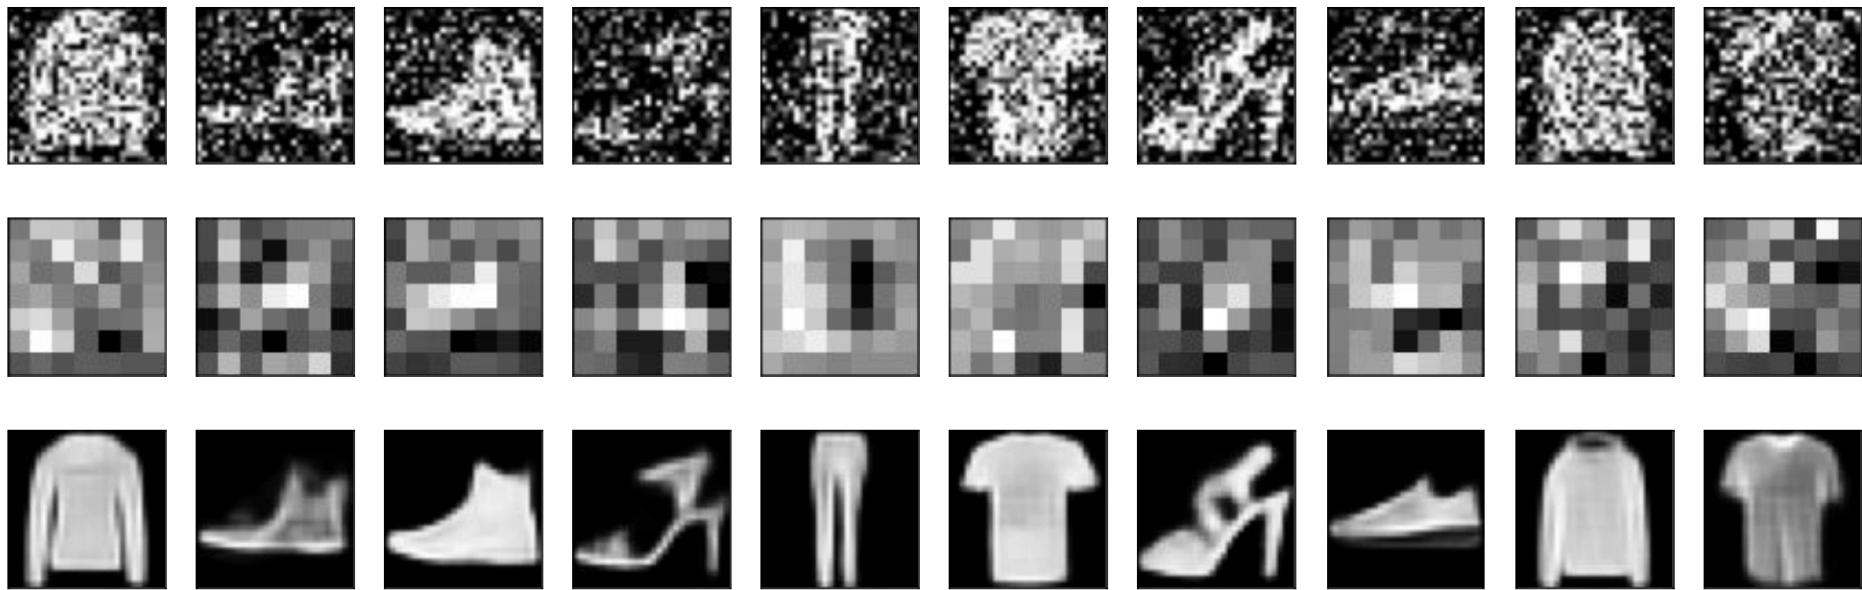
    model = tf.keras.Model(inputs =inputs, outputs=decoder_output)
    encoder_model = tf.keras.Model(inputs=inputs, outputs=encoder_visualization)

    return model, encoder_model
```

Convolutional Auto-Encoders



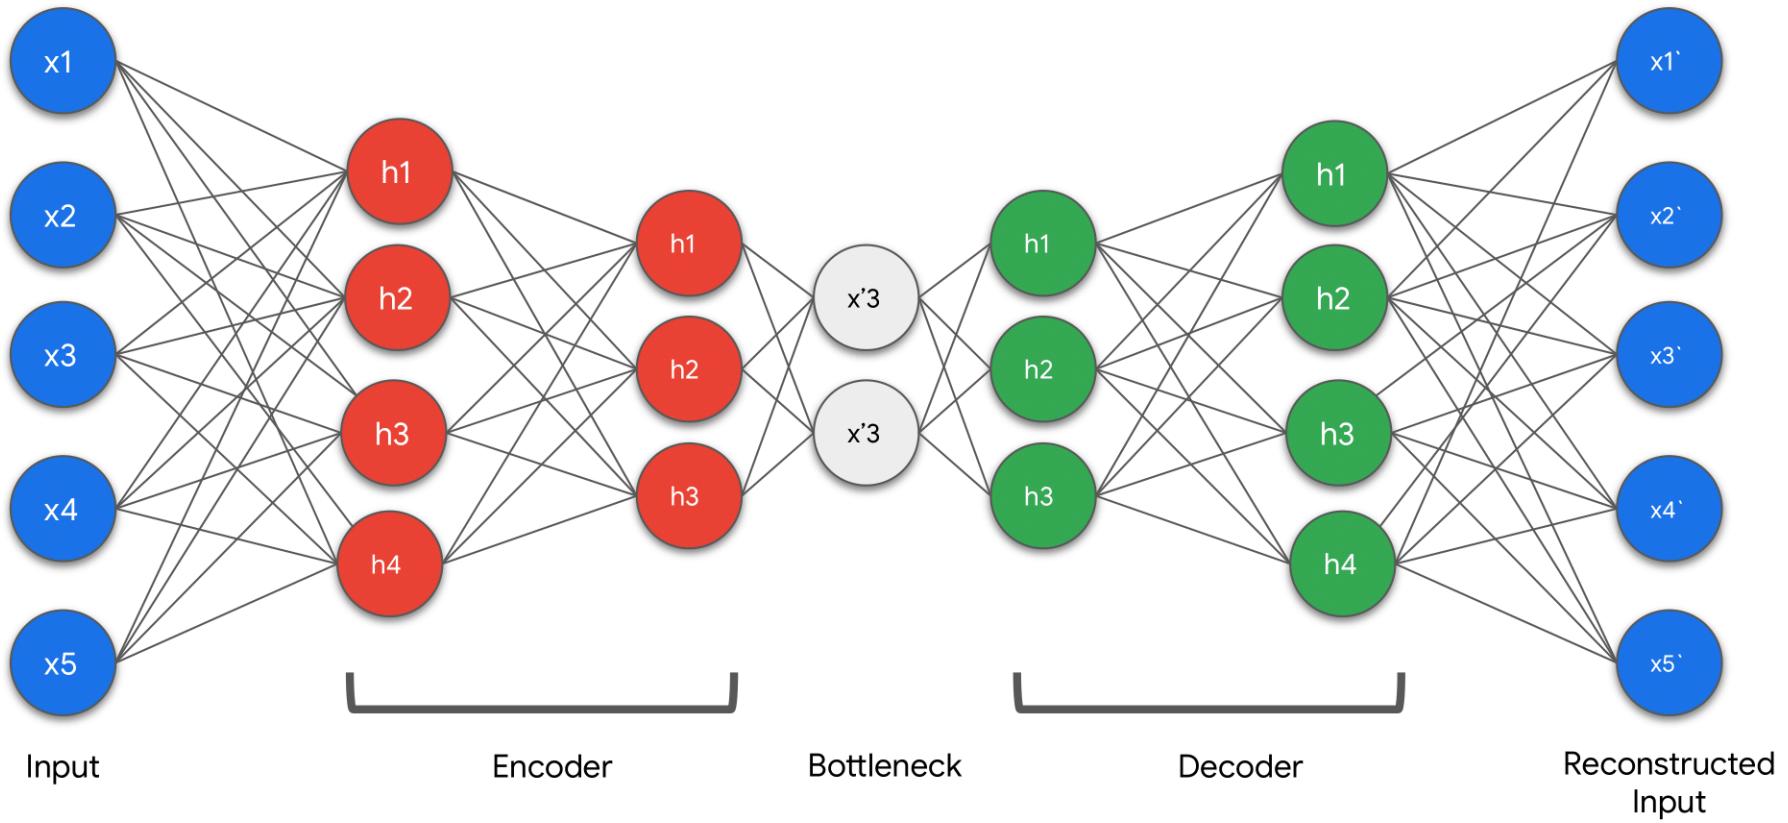


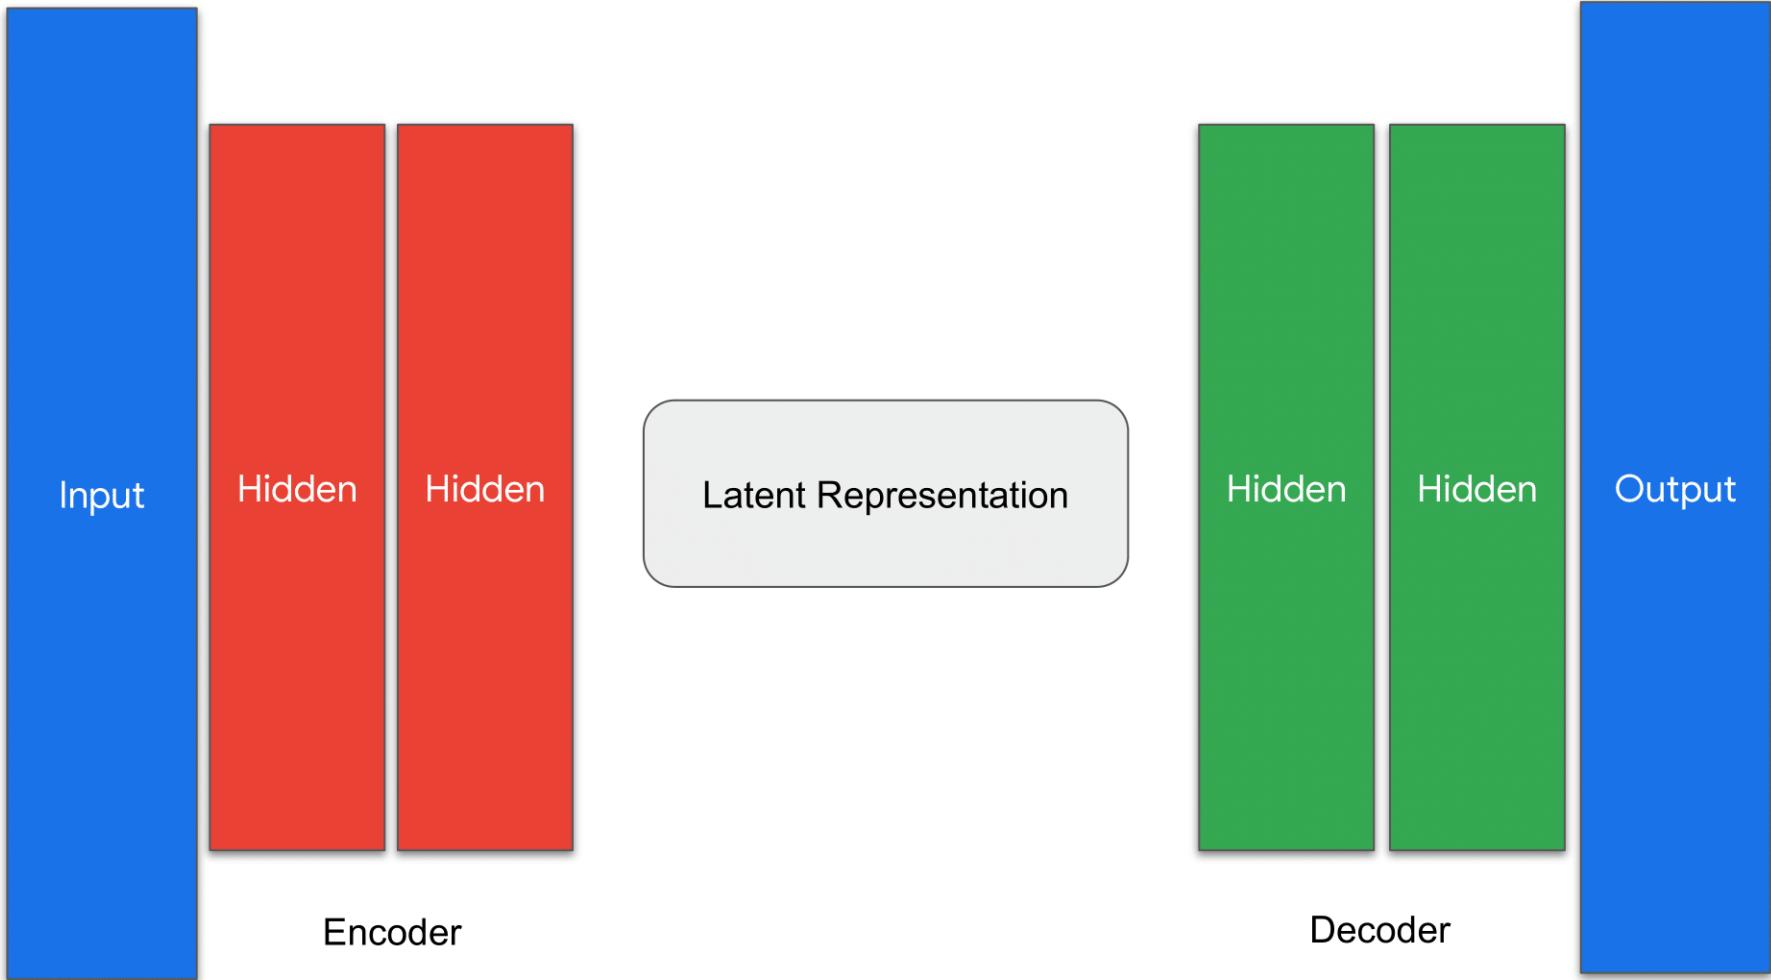


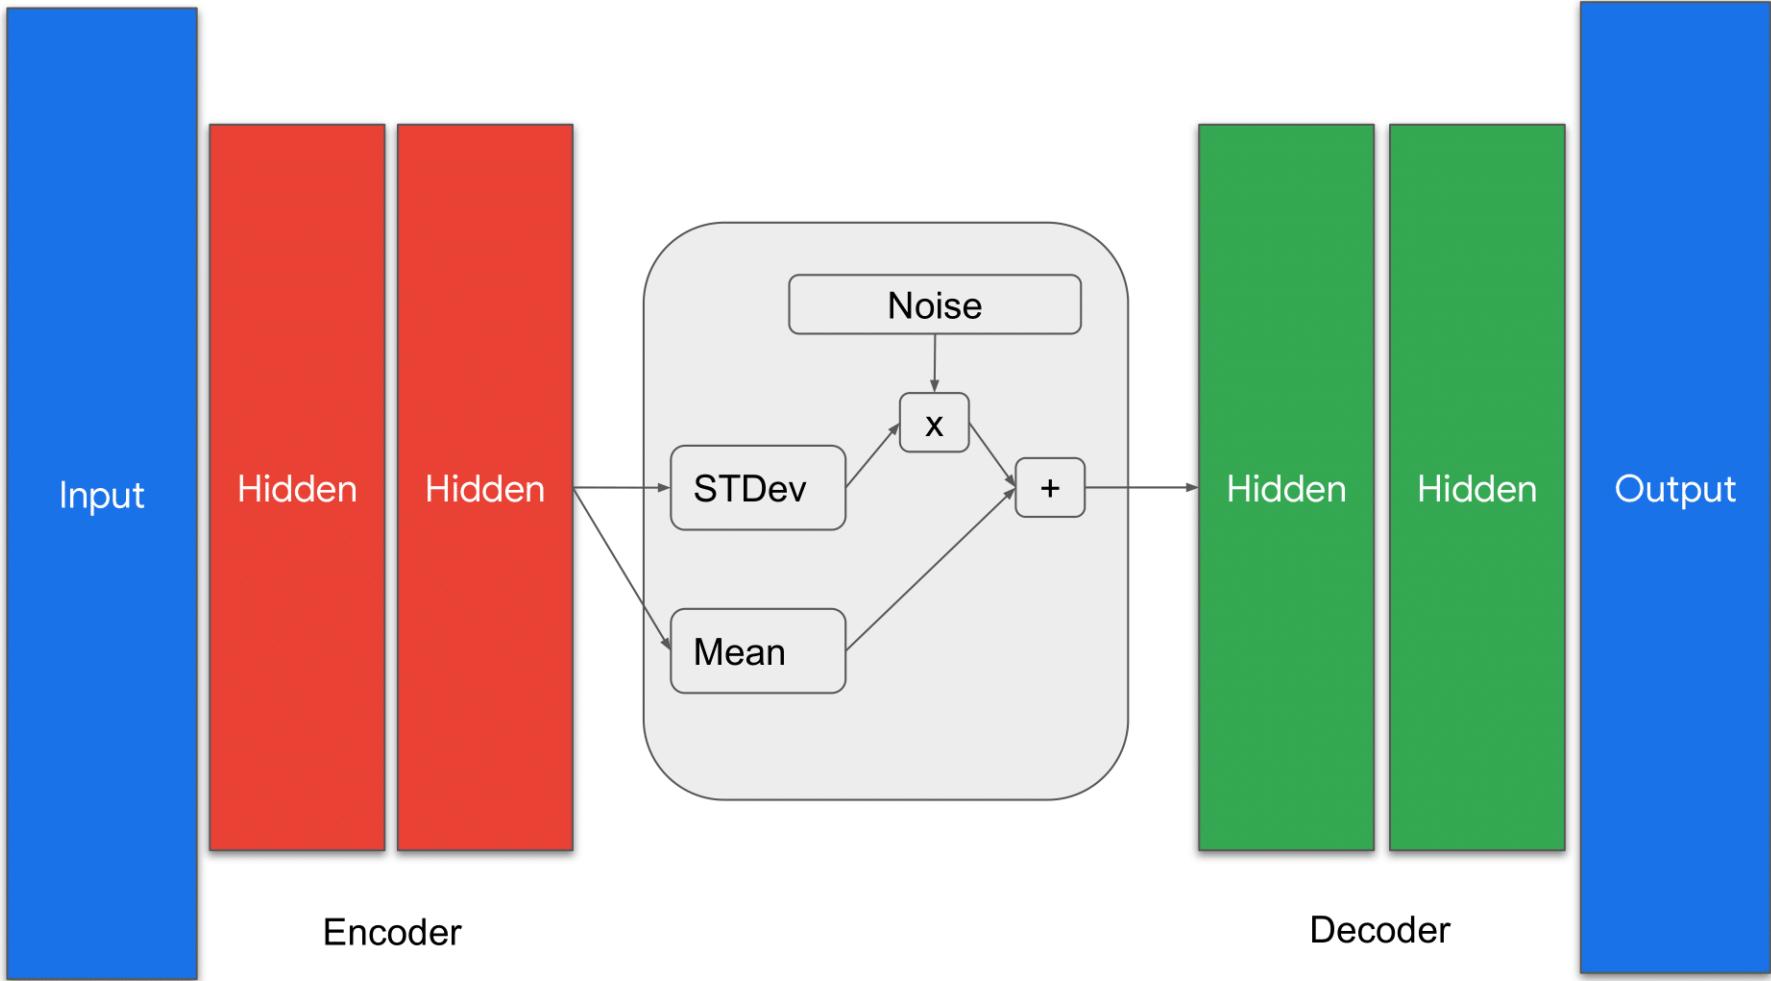
```
def map_image_with_noise(image, label):
    noise_factor = 0.5
    image = tf.cast(image, dtype=tf.float32)
    image = image / 255.0

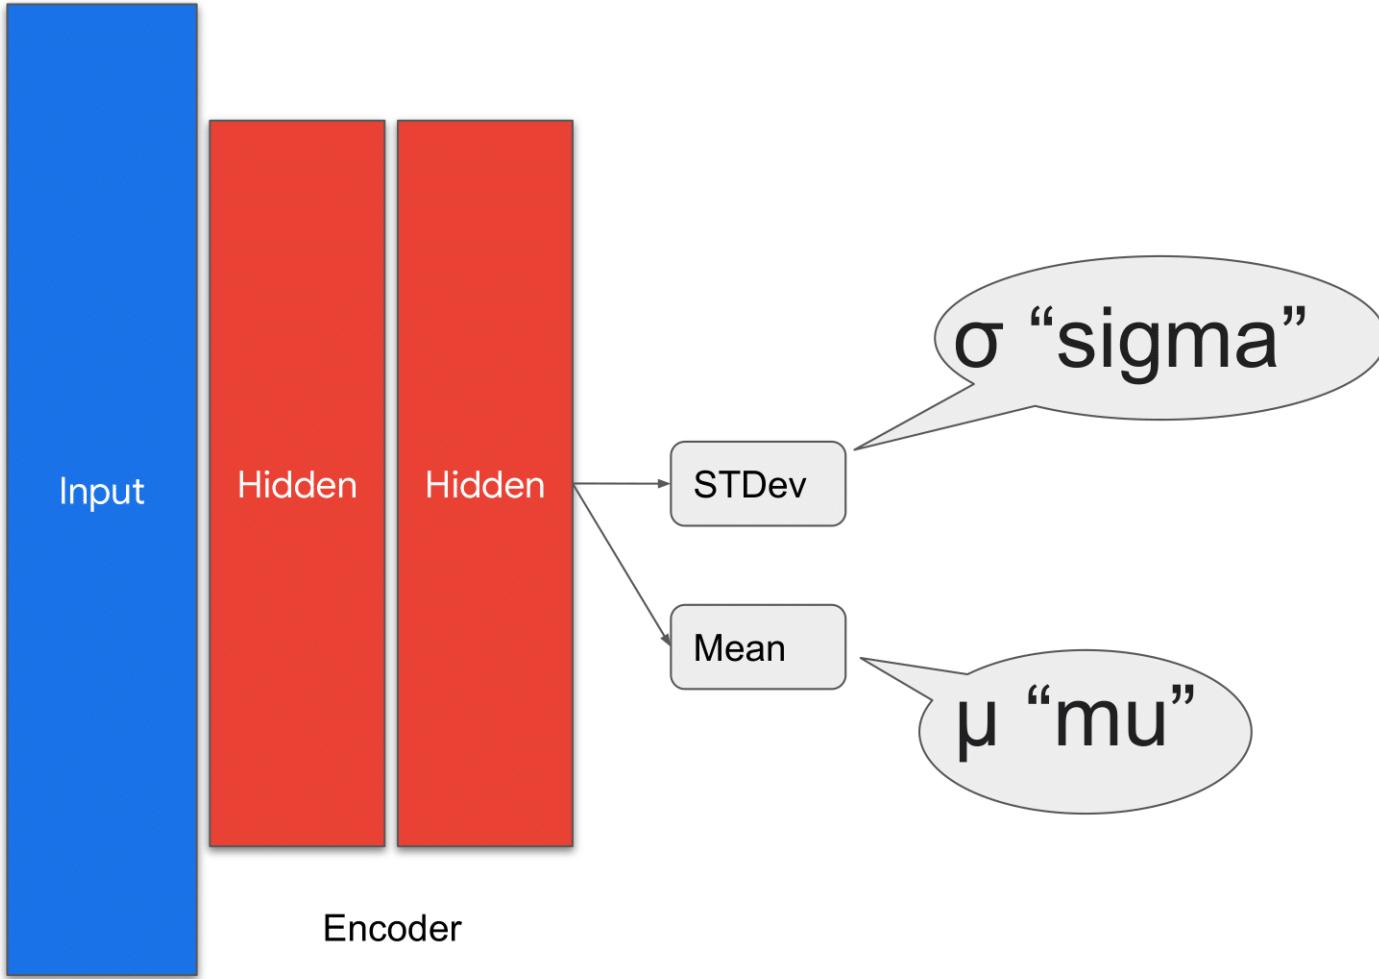
    factor = noise_factor * tf.random.normal(shape=image.shape)
    image_noisy = image + factor
    image_noisy = tf.clip_by_value(image_noisy, 0.0, 1.0)

    return image_noisy, image
```









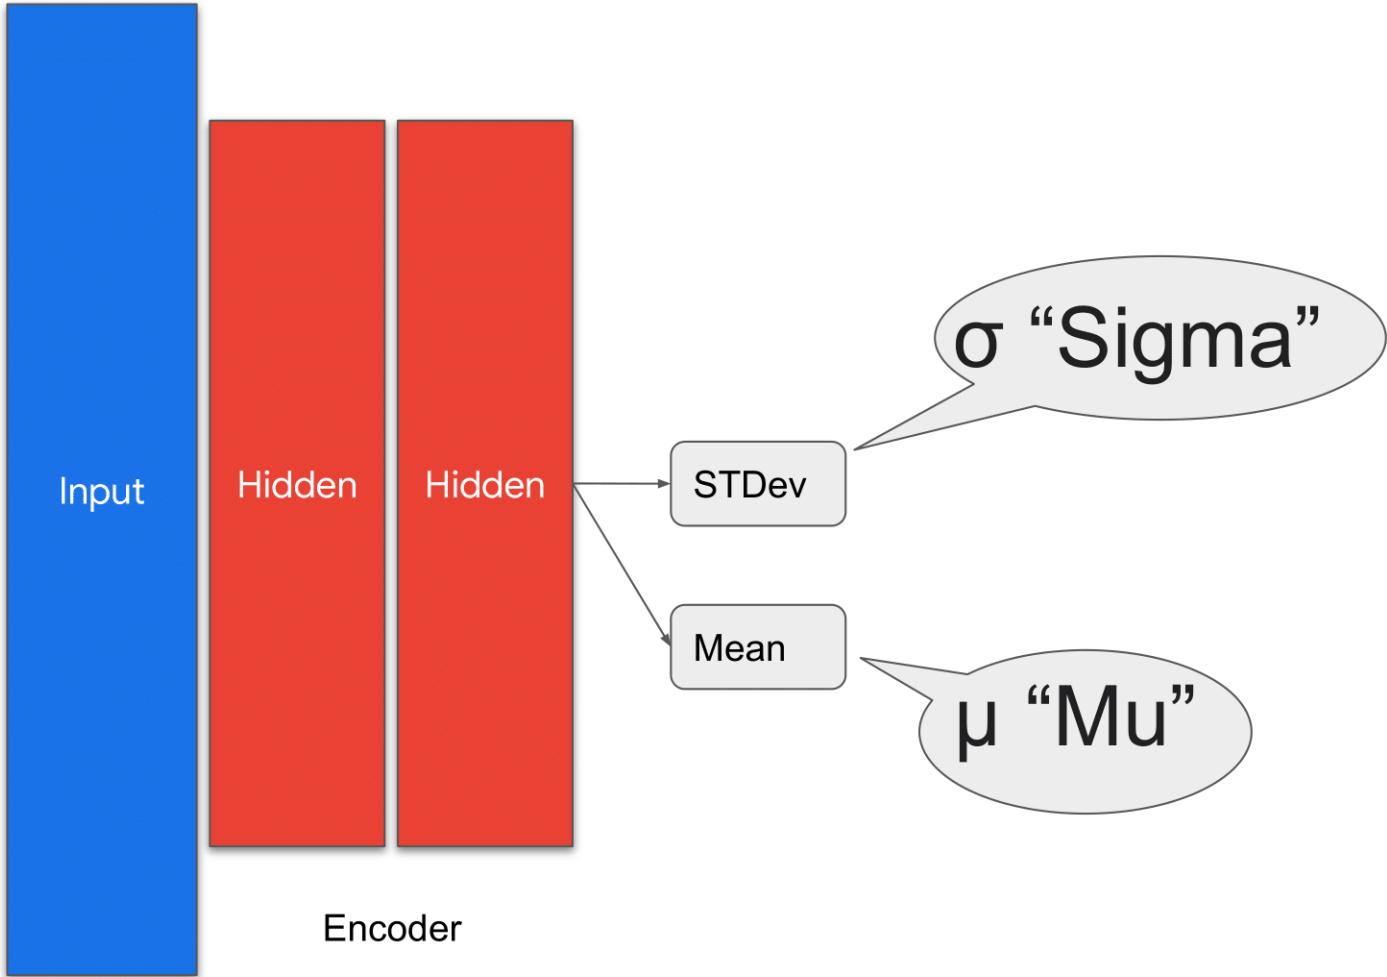
Probability Distribution

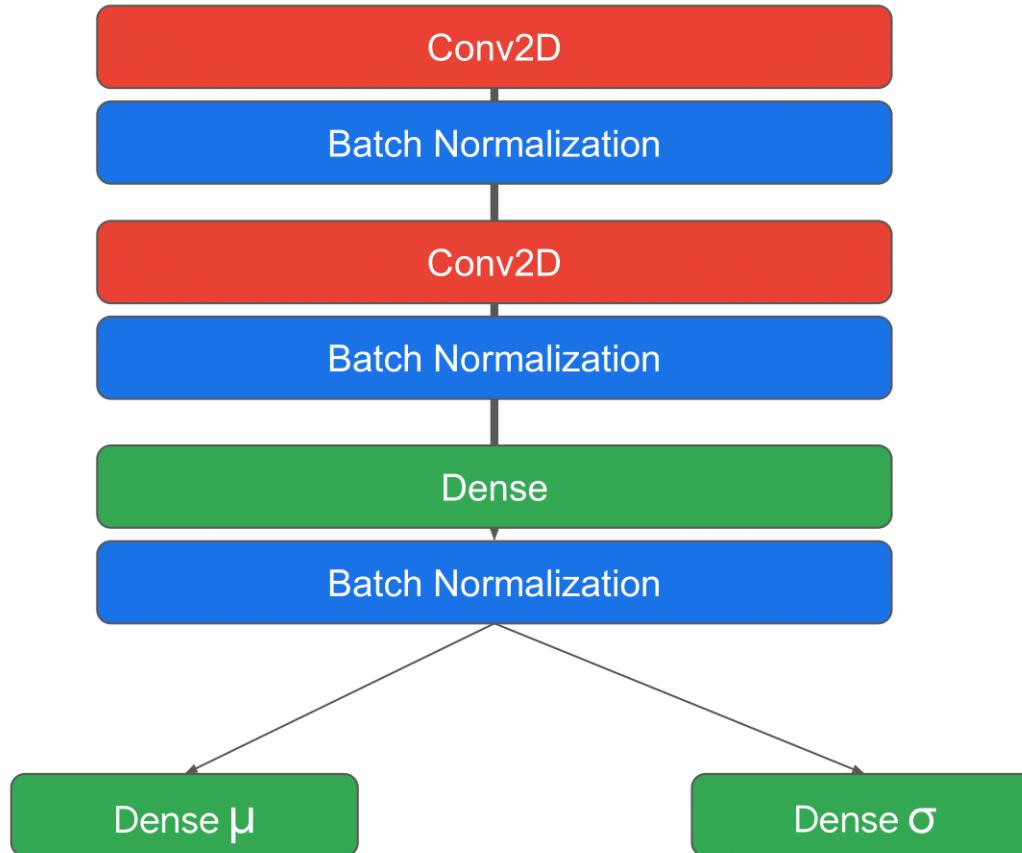
Gaussian probability density function or Normal Distribution.

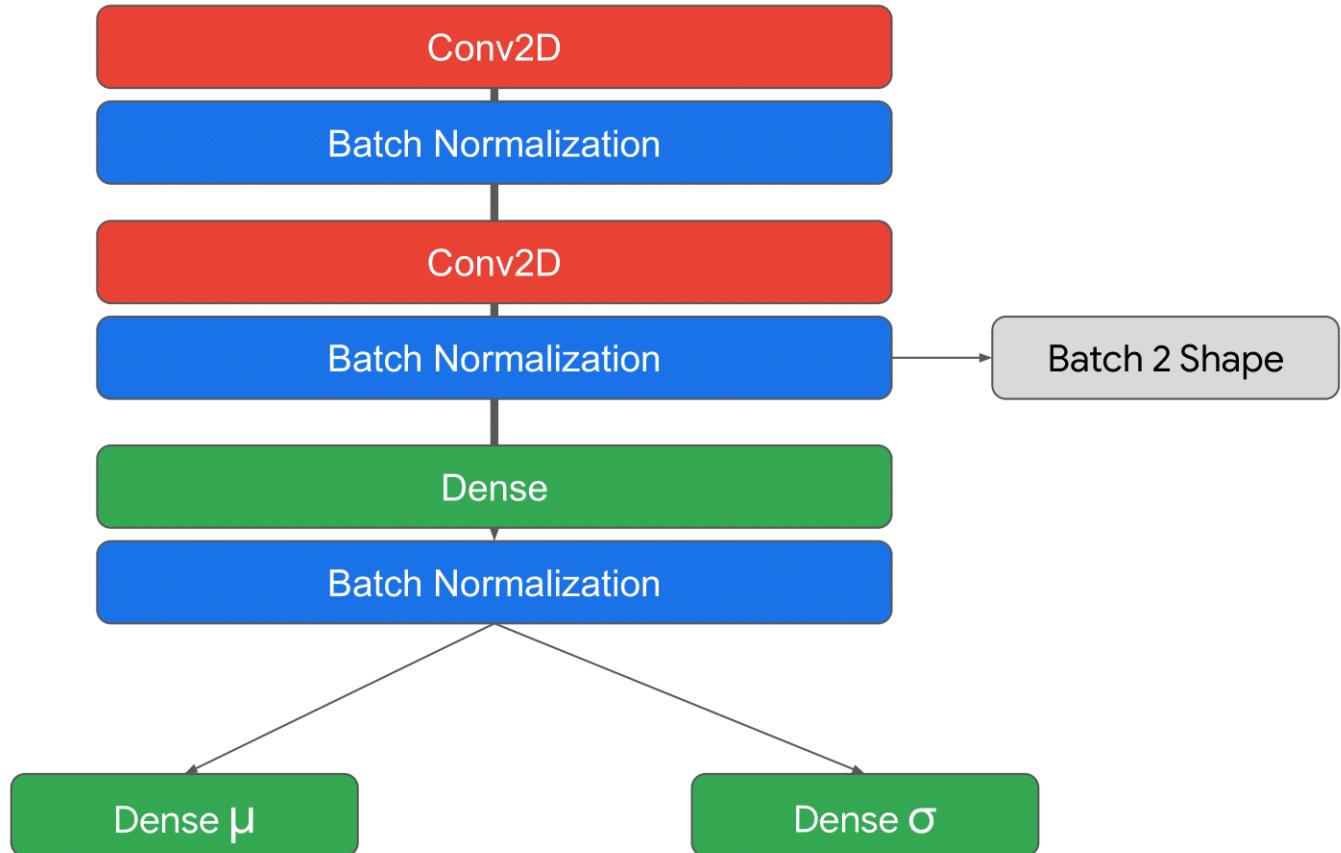
Normal Distribution is controlled by:

- μ “mean”
- σ “standard deviation”

$$N(\mu, \sigma)$$







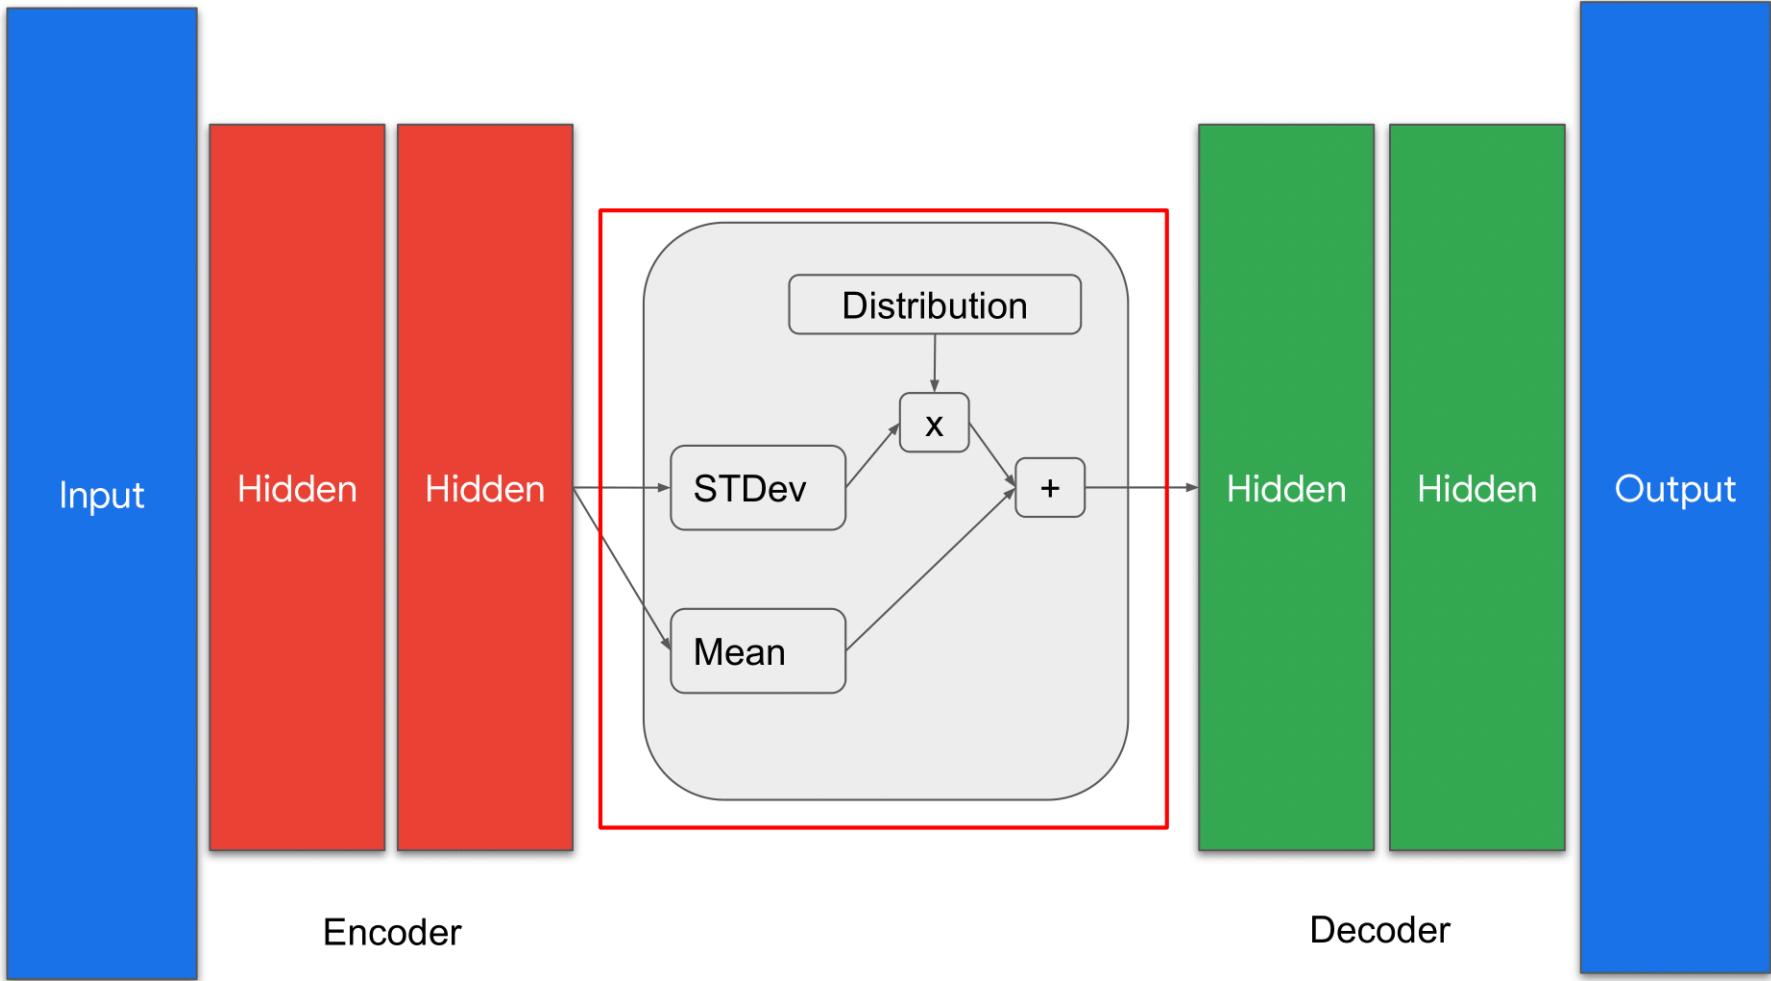
```
# This function defines the encoder's layers
def encoder_layers(inputs, latent_dim):
    x = tf.keras.layers.Conv2D(filters=32, kernel_size=3, strides=2,
                              padding="same", activation='relu',
                              name="encode_conv1")(inputs)
    x = tf.keras.layers.BatchNormalization()(x)

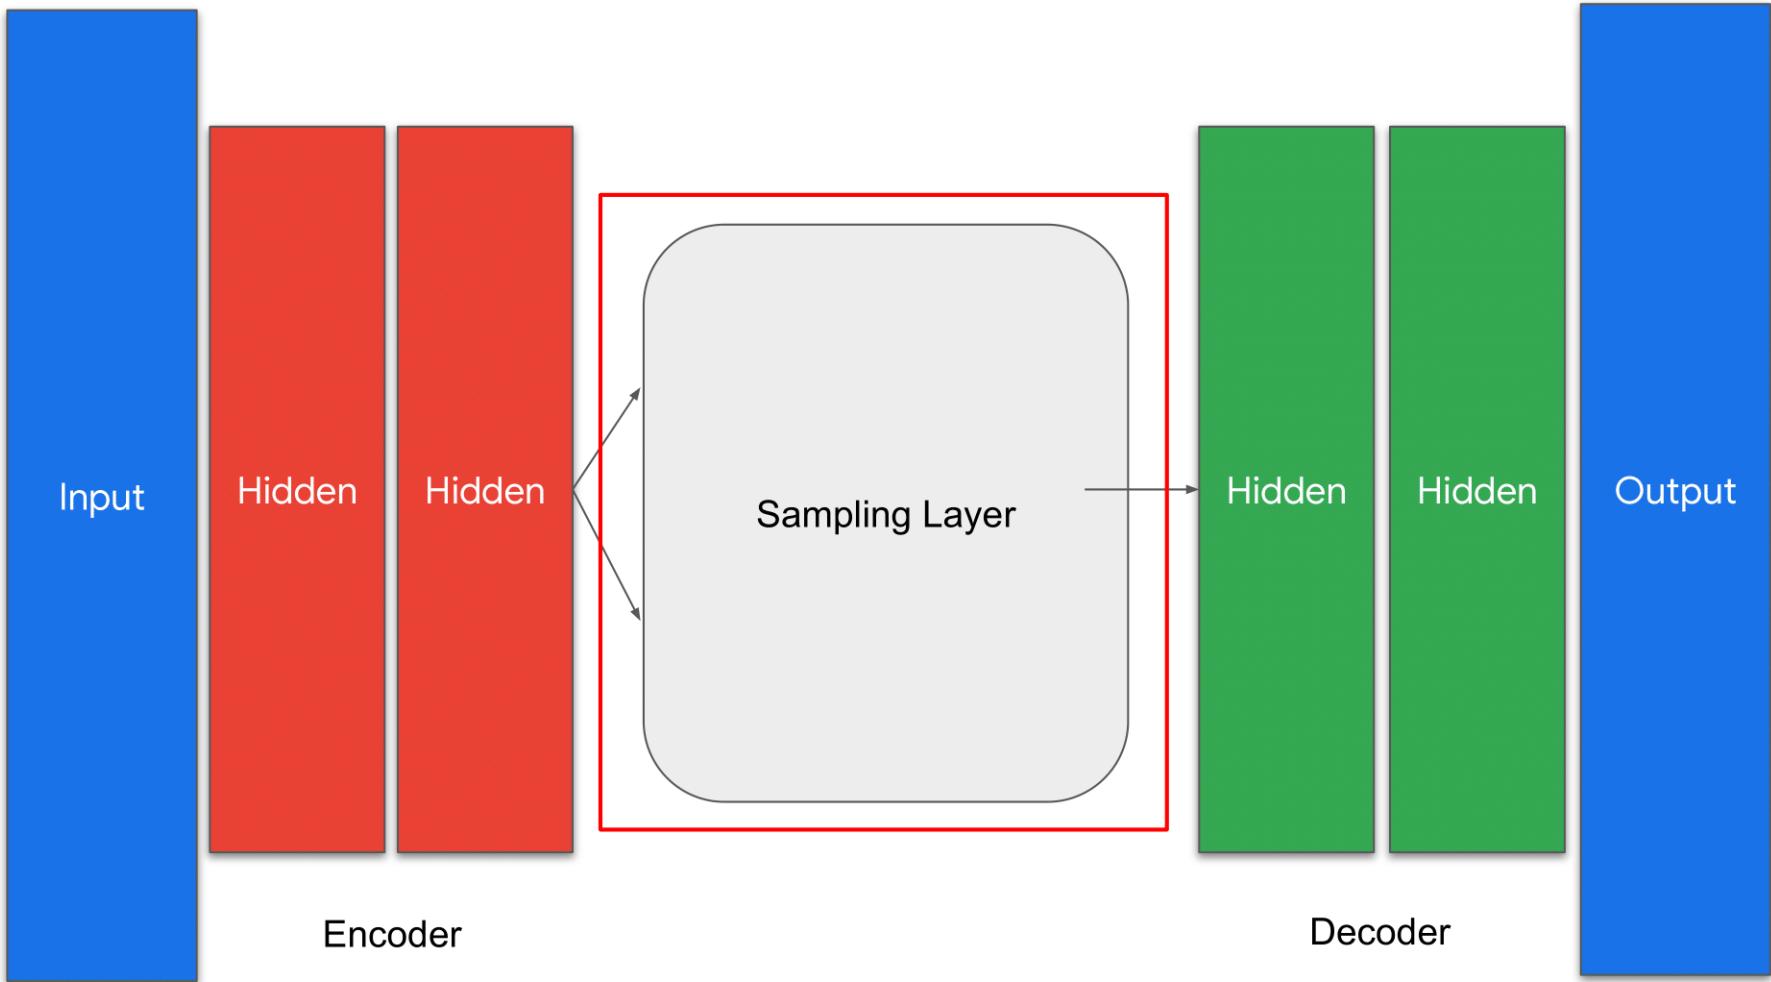
    x = tf.keras.layers.Conv2D(filters=64, kernel_size=3, strides=2,
                              padding='same', activation='relu',
                              name="encode_conv2")(x)
    batch_2 = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.Flatten(name="encode_flatten")(batch_2)
    x = tf.keras.layers.Dense(20, activation='relu', name="encode_dense")(x)
    x = tf.keras.layers.BatchNormalization()(x)

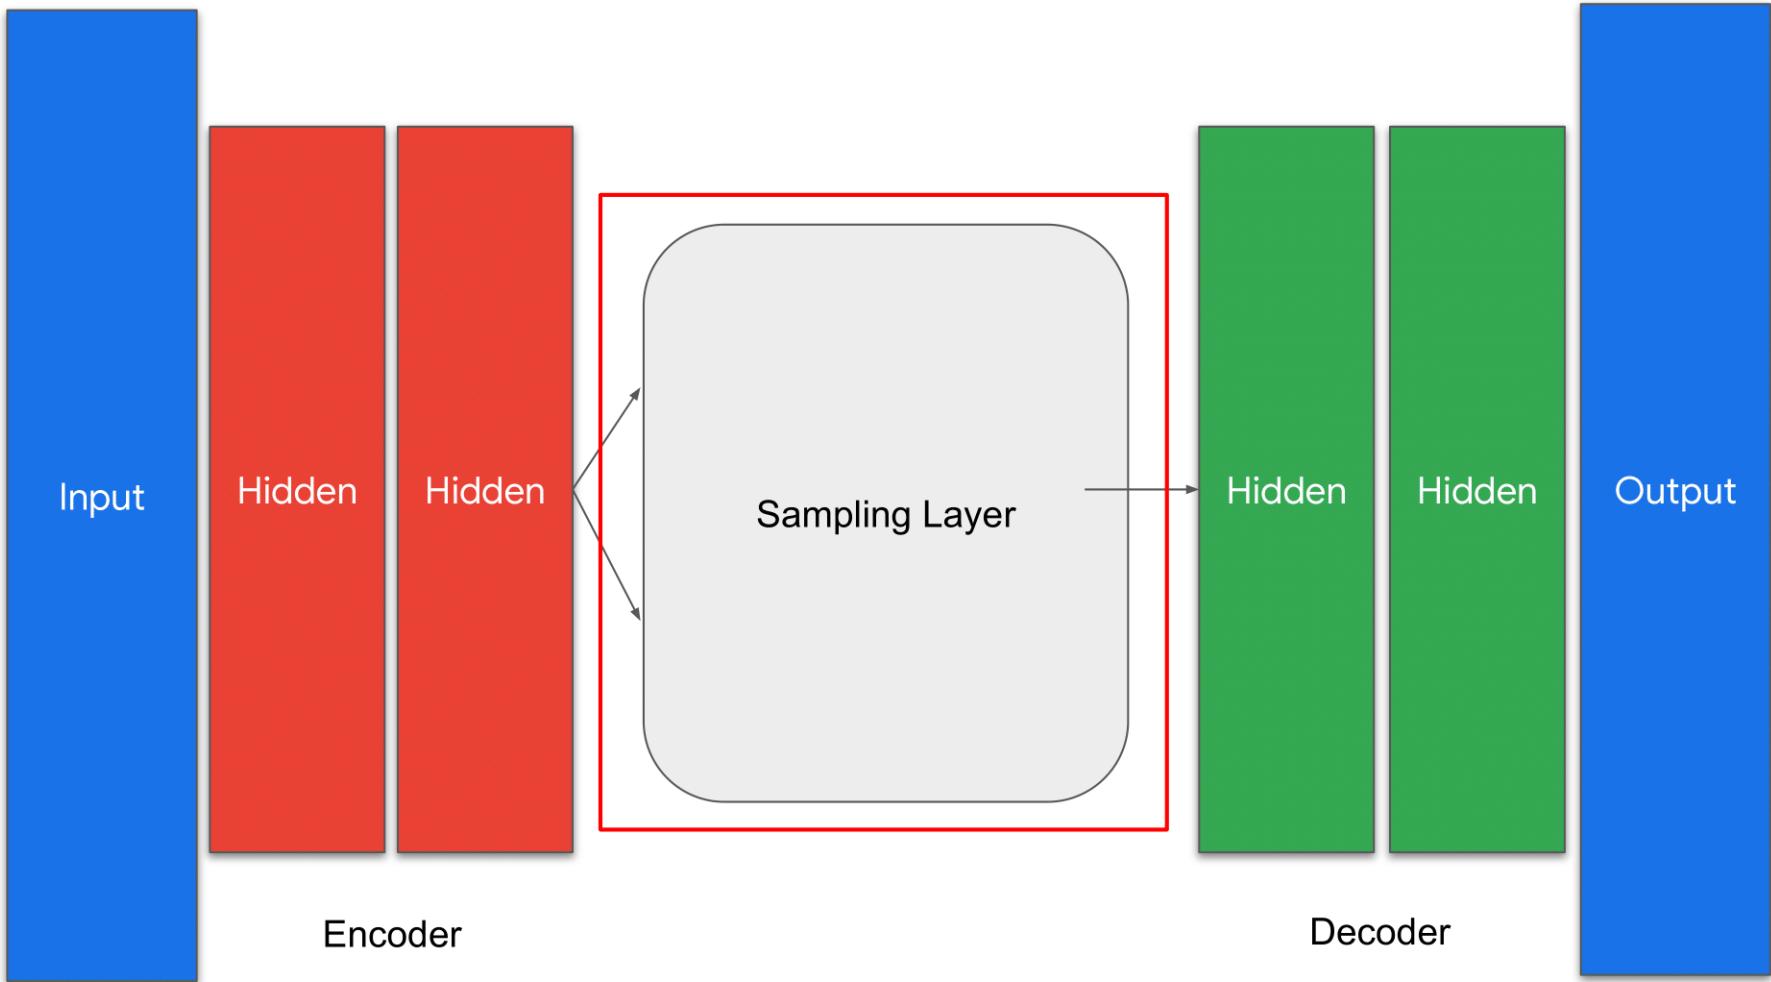
    mu = tf.keras.layers.Dense(latent_dim, name='latent_mu')(x)
    sigma = tf.keras.layers.Dense(latent_dim, name ='latent_sigma')(x)

return mu, sigma, batch_2.shape
```





```
class Sampling(tf.keras.layers.Layer):
    def call(self, inputs):
        mu, sigma = inputs
        batch = tf.shape(mu)[0]
        dim = tf.shape(mu)[1]
        epsilon = tf.keras.backend.random_normal(shape=(batch, dim))
        return mu + tf.exp(0.5 * sigma) * epsilon
```



Input

Hidden

Hidden

Sampling Layer

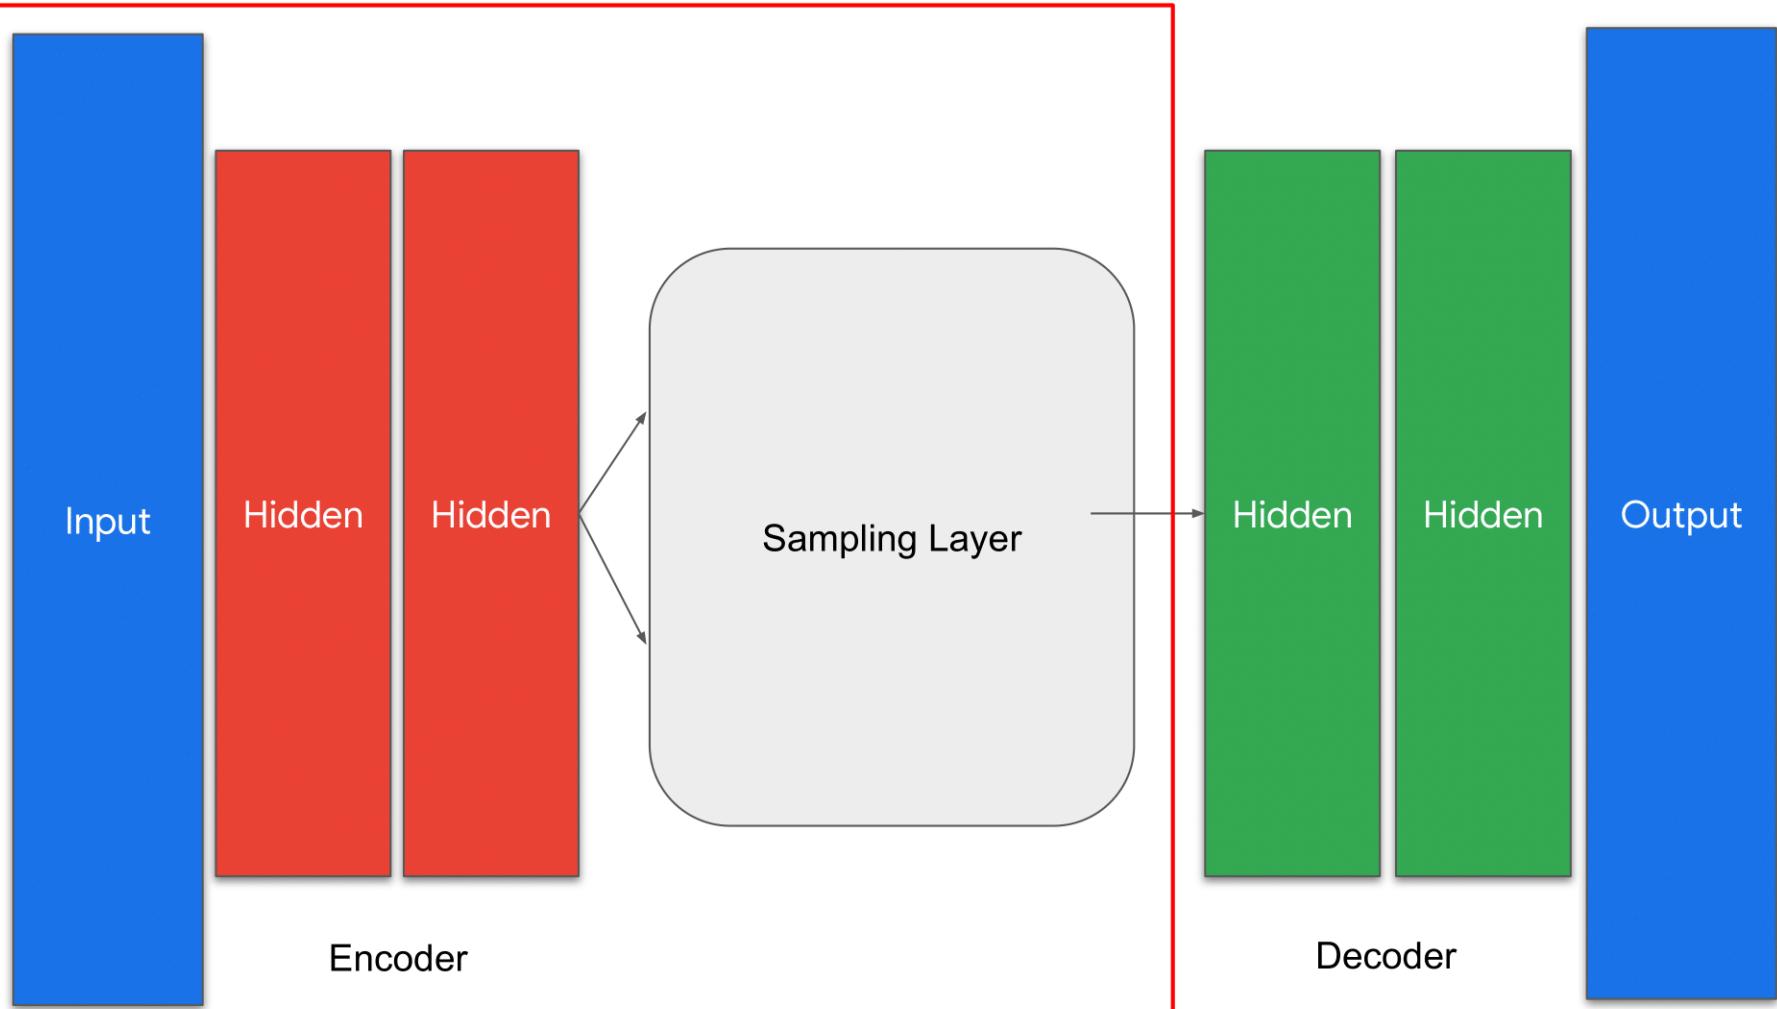
Hidden

Hidden

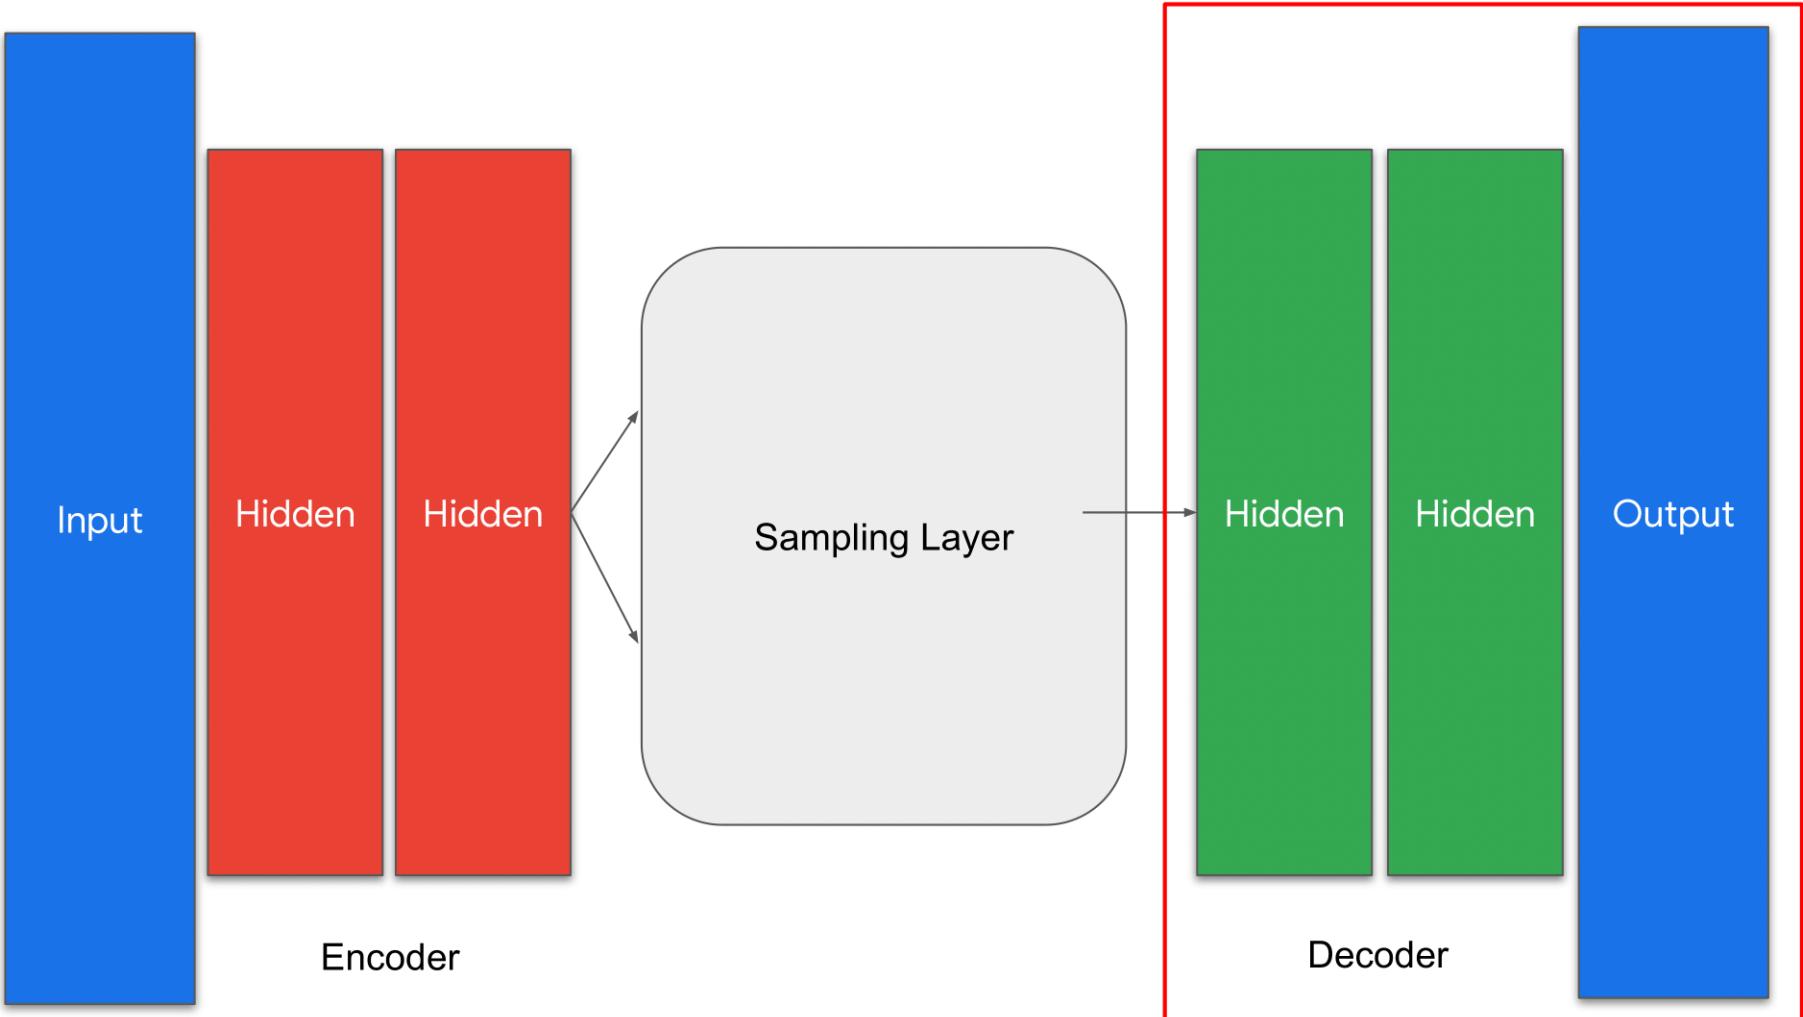
Output

Encoder

Decoder



```
def encoder_model(LATENT_DIM, input_shape):
    inputs = tf.keras.layers.Input(shape=input_shape)
    mu, sigma, conv_shape = encoder_layers(inputs, latent_dim=LATENT_DIM)
    z = Sampling()((mu, sigma))
    model = tf.keras.Model(inputs, outputs=[mu, sigma, z])
    return model, conv_shape
```



```
def decoder_layers(inputs, conv_shape):
    units = conv_shape[1] * conv_shape[2] * conv_shape[3]
    x = tf.keras.layers.Dense(units, activation = 'relu',
                             name="decode_dense1")(inputs)
    x = tf.keras.layers.BatchNormalization()(x)

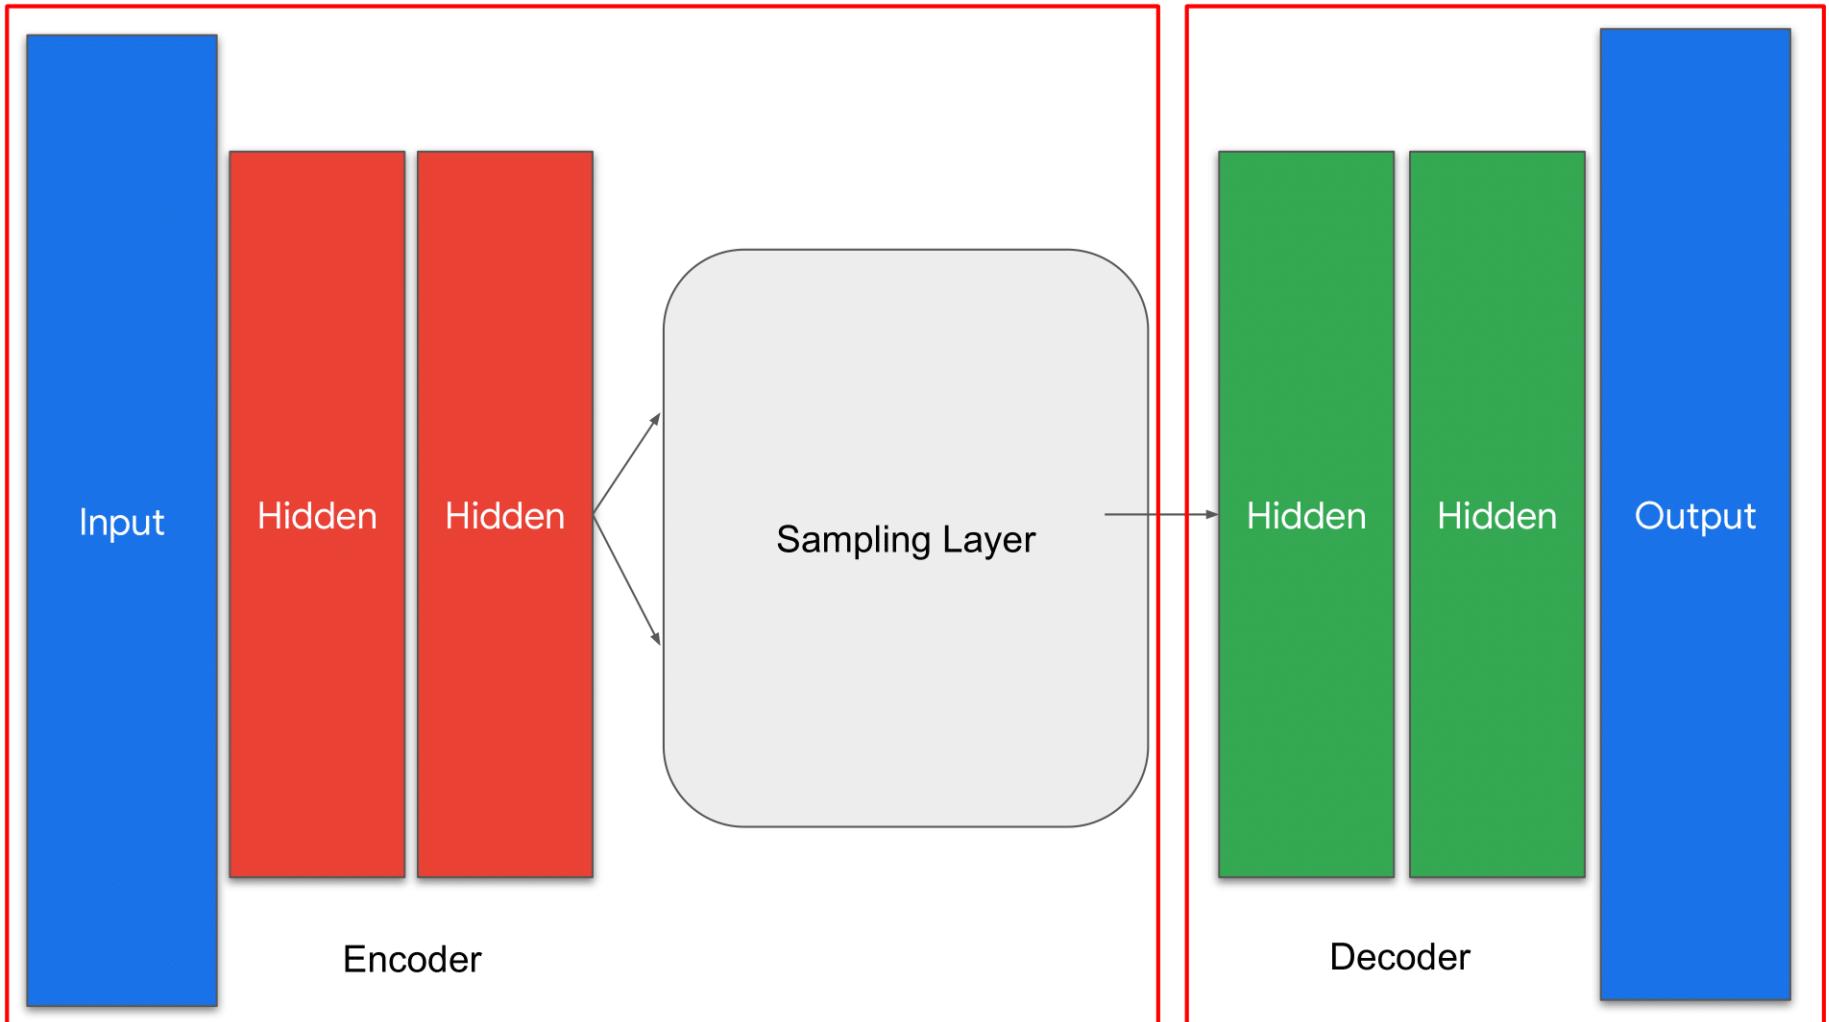
    x = tf.keras.layers.Reshape((conv_shape[1], conv_shape[2], conv_shape[3]),
                                name="decode_reshape")(x)
    x = tf.keras.layers.Conv2DTranspose(filters=64, kernel_size=3, strides=2,
                                       padding='same', activation='relu',
                                       name="decode_conv2d_2")(x)
    x = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.Conv2DTranspose(filters=32, kernel_size=3, strides=2,
                                       padding='same', activation='relu',
                                       name="decode_conv2d3")(x)
    x = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.Conv2DTranspose(filters=1, kernel_size=3, strides=1, padding='same',
                                       activation='sigmoid', name="decode_final")(x)

return x
```

```
def decoder_model(latent_dim, conv_shape):
    inputs = tf.keras.layers.Input(shape=(latent_dim,))
    outputs = decoder_layers(inputs, conv_shape)
    model = tf.keras.Model(inputs, outputs)
    return model
```



```
# Define a kl reconstruction loss function

def kl_reconstruction_loss(inputs, outputs, mu, sigma):
    kl_loss = 1 + sigma - tf.square(mu) - tf.math.exp(sigma)
    return tf.reduce_mean(kl_loss) * -0.5
```

```
def vae_model(encoder, decoder, input_shape):
    inputs = tf.keras.layers.Input(shape=input_shape)
    mu = encoder(inputs)[0]
    sigma = encoder(inputs)[1]
    z = encoder(inputs)[2]
    reconstructed = decoder(z)
    model = tf.keras.Model(inputs=inputs, outputs=reconstructed)
    loss = kl_reconstruction_loss(inputs, z, mu, sigma)
    model.add_loss(loss)
    return model
```

```
for epoch in range(epochs):
    for step, x_batch_train in enumerate(train_dataset):
        with tf.GradientTape() as tape:
            reconstructed = vae(x_batch_train)
            flattened_inputs = tf.reshape(x_batch_train, shape=[-1])
            flattened_outputs = tf.reshape(reconstructed, shape=[-1])
            loss = bce_loss(flattened_inputs, flattened_outputs) * 784
            loss += sum(vae.losses) # Add KLD regularization loss

        grads = tape.gradient(loss, vae.trainable_weights)
        optimizer.apply_gradients(zip(grads, vae.trainable_weights))
```

epoch: 99, step: 400

